MODULE *PrimSecRep*

EXTENDS *Util*, *TLC*

CONSTANTS Values:

Rep, replicas

Val,

**HOLES*:

Types:

Object, object *IDs* or addresses

DataType, data type

WriteType, write structure (the one that gets written to the data)

ErrorTypes, different errors returned by the secondaries

Abstract methods ("holes") :

StoreUpdate(_, _, _), (*r*,*o*,*w*) perform write *w* (∈ *WriteType*) on local data of object *o* of replica *r*

StoreRead(_, _), *read*(*r*, *o*) peform local read at replica *r* of object *r*. Returns the value

PrepareWr(_, _, _), *prepareWr*(*r*, *o*, *w*) prepares write *w* to send out to secondaries. This preparation is done by the primary.

ReplyFinishedWrite(_, _, _), (*r*, *o*, *w*) replica *r* replies to a finished write on object *o*

GetWriteResponse(_, _, _, _, _), (*r*,*o*,version,sender) returns the response that should be sent out In case of some error (*e.g.*, wrong vers), should return [*type* ↦ "*noResp*"], in which case nothing will be committed to disk or sent out

Reply(_, _, _), (*r*, type, reply) reply to client

NoReply, no reply is sent yet, so just keep the aux variables unchanged

InitDataVal, initial data value for the disks

InitPrim(_), *InitPrim*(*o*) = value of the initial primary of object *o*

InitSec(_), initial secondaries for object *o*

NoRep,

Client,

NoWrite

VARIABLES *master*, reliable master state

data, data at each replica

cache, cache of the master state at each replica

stat, state of each replica + object locks of each replica

channel, in-channel of each replica

resps write acknowledgements for each replica from the others

$vars \triangleq \langle master, cache, data, channel, resps, stat \rangle$

$NoRep \triangleq$ CHOOSE $v : v \notin Rep$
$Client \triangleq$ CHOOSE $v : v \notin Rep \land v \neq NoRep$
$NoWrite \triangleq$ CHOOSE $v : v \notin WriteTypes$

$WrReqMsg \triangleq$ Write requests
$[type : \{\text{“wrReq”}\},\ sender : Rep \cup \{Client\},\ object : Object,$
$w : WriteType,\ version : Nat]$
$WrRespMsg \triangleq$ Write responses
$[type : \{\text{“wrResp”}\},\ sender : Rep,\ object : Object,\ ack : (\{\text{“ok”}\} \cup ErrorTypes),$
$w : WriteType,\ version : Nat]$
$Messages \triangleq$ Message types
$WrReqMsg \cup WrRespMsg$

$TypeInvariant \triangleq$
$\land\ Print\ (\text{“1”}, \text{TRUE})$
$\land\ master \in$
$[objects : [Object \to [version : Nat,\ prim : Rep \cup \{NoRep\},\ sec : \text{SUBSET}\ (Rep)]],$
$health : [Rep \to \{\text{“dead”},\ \text{“alive”}\}]]$

$\land\ Print\ (\text{“2”}, \text{TRUE})$
$\land\ data \in [Rep \to [Object \to DataType]]$
$\land\ Print\ (\text{“3”}, \text{TRUE})$
$\land\ cache \in [Rep \to [Object \to [version : Nat,\ prim : Rep \cup \{NoRep\},\ sec : \text{SUBSET}\ (Rep)]]]$
$\land\ Print\ (\text{“5”}, \text{TRUE})$
$\land\ channel \in [Rep \to Seq(Messages)]$

$\land\ Print\ (\text{“7”}, \text{TRUE})$
$\land\ stat \in [Rep \to [phase : \{\text{“alive”},\ \text{“dead”}\},\ lock : [Object \to \{\text{“rdy”},\ \text{“busy”}\}],$
$in\_progress : [Object \to [val : WriteType,\ version : Nat]]]]$

$\land\ Print\ (\text{“8”}, \text{TRUE})$
$\land\ resps \in [Rep \to [Object \to [Rep \to \{\text{“ok”},\ \text{“waiting”},\ \text{“timeout”},\ \text{“n/a”}\} \cup ErrorTypes]]]$
separate set of responses per object
$\land\ Print\ (\text{“9”}, \text{TRUE})$
$Init \triangleq$
$\land\ master =$
$[objects \mapsto [o \in Object \mapsto$
$[version \mapsto 0,\ prim \mapsto InitPrim(o),\ sec \mapsto InitSec(o)]],$
$health \mapsto [r \in Rep \mapsto \text{“alive”}]]$ but all replicas are alive

$\land\ cache = [r \in Rep \mapsto [o \in Object \mapsto$
$[version \mapsto 0,\ prim \mapsto InitPrim(o),\ sec \mapsto InitSec(o)]]]$

$\land\ data = [r \in Rep \mapsto [o \in Object \mapsto InitDataVal]]$
$\land\ channel = [r \in Rep \mapsto \langle\rangle]$ no message yet
$\land\ resps = [r \in Rep \mapsto [o \in Object \mapsto [r1 \in Rep \mapsto \text{“waiting”}]]]$

$\land\ stat = [r \in Rep \mapsto [phase \mapsto$ "alive", $lock \mapsto [o \in Object \mapsto$ "rdy"],
$in\_progress \mapsto [o \in Object \mapsto [val \mapsto NoWrite,\ version \mapsto 0]]]]$
all reps alive and ready to accept updates

MASTER functioning

In the most general case, the master does not participate actively. It's only for mitigation and for a reliable source of accurate information. In master-monitored replicated systems (like *GFS*), the master is active and monitors replicas, releases primary leases, removes replicas, adds replicas, etc. In self-monitoring replica sys, (like Blue), master is not active in the spec.

$MasterActions \triangleq$ FALSE

TIME AND CHANNELS: (sources of timeouts and losses)

The channel to replica r loses one of the messages that should have arrived there.

$\_TransLoss(r) \triangleq$
$\land\ channel[r] \neq \langle\rangle$ the channel is not empty
$\land\ channel' = [channel$ EXCEPT $![r] = Tail(@)]$
$\land$ UNCHANGED $\langle master,\ cache,\ data,\ stat,\ resps\rangle$
$TransLoss \triangleq \exists\, r \in Rep : \_TransLoss(r)$

(Ex-)Primary r times out waiting for response from replica s.

$\_Timeout(r,\ o,\ s) \triangleq$
LET $updateMsgIsLost(chan,\ update) \triangleq$
$\forall\, i \in 1\,..\,Len(chan) : chan[i].w \neq update.val$
IN
$\land\ stat[r].phase =$ "alive" r is alive
$\land\ stat[r].lock[o] =$ "busy" r is indeed waiting for responses to a write request
$\land\ resps[r][o][s] =$ "waiting" r is still waiting for response from s
$\land\ updateMsgIsLost(channel[s],\ stat[r].in\_progress[o])$ * REDUCE SPACE: the update message sent to s was lost – *TODO*
$\land\ resps' = [resps$ EXCEPT $![r][o][s] =$ "timeout"] response of s timed out
$\land$ UNCHANGED $\langle master,\ cache,\ data,\ stat,\ channel\rangle$
$Timeout \triangleq$
$\exists\, r \in Rep,\ o \in Object,\ s \in Rep :$
$\land\ r \neq s$
$\land\ \_Timeout(r,\ o,\ s)$

$TimeActions \triangleq$
$\lor\ TransLoss \qquad \lor\ Timeout$

REPLICA actions.

Replica r suddenly dies.

$\_ReplicaDeath(r) \triangleq$
$\quad \land\ stat[r].phase = \text{“alive”}$ replica used to be alive

$\quad \land\ stat' = [stat \text{ EXCEPT } ![r].phase = \text{“dead”}]$ declares itself dead
$\quad \land \text{ UNCHANGED } \langle master,\ cache,\ data,\ channel,\ resps \rangle$
$ReplicaDeath \triangleq \exists\, r \in Rep : \_ReplicaDeath(r)$

Replica r updates its cache $w/$ the accurate version from master. One could imagine multiple times when this could be triggered in the actual implem of, $e.g.$, GFS ($e.g.$, during $HeartBeat$ protocol, after a client comes $w/$ a higher $version \neq$, etc.). Since I don't know what happens in the real protocol exactly, I will assume it can happen anytime.

$\_ReadVersion(r,\ o) \triangleq$
$\quad \land\ stat[r].phase \neq \text{“dead”}$ replica is not dead yet
$\quad \land\ cache[r][o].version \neq master.objects[o].version$ r's cache is out-of-date

$\quad \land\ cache' = [cache \text{ EXCEPT } ![r][o] = master.objects[o]]$ cache new version
$\quad \land \text{ UNCHANGED } \langle master,\ data,\ channel,\ resps,\ stat \rangle$
$ReadVersion \triangleq$
$\quad \land \exists\, r \in Rep,\ o \in Object : \_ReadVersion(r,\ o)$

Replica r drops a write req w. The reasons might be various. ($E.g.$, in GFS, one reason can be that prim cannot find the data associated $w/$ this request in its LRU).

$DropWrite(r,\ o,\ w) \triangleq$ changes $\langle data,\ channel,\ stat,\ resps \rangle$
$\quad \land \text{ UNCHANGED } \langle data,\ channel,\ stat,\ resps \rangle$

Primary continues $w/$ a client request. It pushes the data to its local store (always considered to succeed from this at this point, b/c anyway I provide for general request dropping). The primary then sends update messages to all secondaries. The write needs preparation before it's given out to secondaries. In GFS ($e.g.$), the preparation means deciding on an address to write at and setting this adr in the msg.

$PrimaryContinuesWrite(r,\ o,\ w,\ allreplicas,\ version) \triangleq$ changes $\langle data,\ channel,\ stat,\ resps \rangle$, sends $Reply$
$\quad \land \text{ LET}$
$\qquad prepared\_wr \triangleq PrepareWr(r,\ o,\ w)$
$\qquad wr\_req \triangleq [type \mapsto \text{“wrReq”},\ sender \mapsto r,\ object \mapsto o,$
$\qquad\qquad version \mapsto version,\ w \mapsto prepared\_wr]$
$\qquad\qquad$ the write needs to be prepared for the secondaries
$\quad \text{IN}$
$\qquad \land \text{ IF } prepared\_wr \neq NoWrite \text{ THEN}$ No error while preparing

$\qquad\quad \land\ StoreUpdate(r,\ o,\ prepared\_wr)$ perform the update locally
$\qquad\quad \land\ channel' = [r1 \in Rep \mapsto$ send the update req to secondaries
$\qquad\qquad \text{IF } r1 \in allreplicas \setminus \{r\} \text{ THEN}$
$\qquad\qquad\quad channel[r1] \circ \langle wr\_req \rangle$
$\qquad\qquad \text{ELSE } channel[r1]]$ not a secondary

$\qquad\quad \land\ stat' = [stat \text{ EXCEPT } ![r].lock[o] = \text{“busy”},$ the replica becomes busy
$\qquad\qquad\qquad\qquad$ and will not accept further updates on object o
$\qquad\qquad\qquad\qquad$ until this write finishes.

$!\![r].in\_progress[o] = [val \mapsto prepared\_wr,\ version \mapsto version]]$
maintain the write, for future reference
$\land\ resps' = [resps$ EXCEPT $!\![r][o] = [r1 \in Rep \mapsto$
IF $r = r1$ THEN "ok" this replica has already answered
ELSE IF $r1 \in allreplicas$
THEN "waiting" wait for *ack* from secondaries
ELSE "n/a"]] don't wait for *ack* from non-secondaries
$\land\ NoReply$ don't return yet the response

$\land$ UNCHANGED *data*

ELSE Preparing the message resulted in error, indicating
that this write shouldn't go ahead.
Skip it and announce the client of error
$\land$ UNCHANGED $\langle data,\ channel,\ stat,\ resps\rangle$
$\land\ Reply(r,$ "NoWrite", $prepared\_wr)$

Primary r handles a write request. Note that the writes are blocking and processed by one primary one at a time. Only after a write finishes, does the primary start another one.

$PrimaryWrite(r,\ o,\ w,\ ver) \triangleq$ changes $\langle data,\ channel,\ stat,\ resps\rangle$, sends *Reply*
$\land\ cache[r][o].prim = r$ r believes itself to be the primary
$\land\ stat[r].lock[o] =$ "rdy" the primary doesn't have other writes in progress on object o

$\land\ \lor \land PrimaryContinuesWrite(r,\ o,\ w,\ \{cache[r][o].prim\} \cup cache[r][o].sec,\ cache[r][o].version)$
EITHER: perform the write

$\lor \land DropWrite(r,\ o,\ w)$ * OR: drop the request alltogether
$\land\ Reply(r,$ "wrFinished", $[ack \mapsto$ "error", $object \mapsto o,\ w \mapsto w])$

$VersionBased\_GetWriteResponse(r,\ o,\ w,\ version,\ sender) \triangleq$
IF $(cache[r][o].version \leq version)$ THEN good version
$[type \mapsto$ "wrResp", $sender \mapsto r,\ object \mapsto o,$
$w \mapsto w,\ ack \mapsto$ "ok", $version \mapsto cache[r][o].version]$
ELSE bad version
$[type \mapsto$ "wrResp", $sender \mapsto r,\ object \mapsto o,$
$w \mapsto w,\ ack \mapsto$ "badver", $version \mapsto cache[r][o].version]$

Secondary continues the write request from prim. It pushes the data to its local store (always considered to succeed from this at this point, b/c anyway I provide for general request dropping & version has been or doesn't need to be checked.

$SecondaryContinuesWrite(r,\ o,\ w,\ ver,\ sender) \triangleq$ changes $\langle data,\ channel\rangle$
LET
$wr\_resp \triangleq GetWriteResponse(r,\ o,\ w,\ ver,\ sender)$
IN
IF $wr\_resp.type =$ "wrResp" THEN
$\land\ wr\_resp.ack =$ "ok" $\Rightarrow StoreUpdate(r,\ o,\ w)$ store the update persistenly locally
$\land\ wr\_resp.ack \neq$ "ok" $\Rightarrow$ UNCHANGED *data* don't store, there's an error

$\land\ channel' = [channel$ EXCEPT
$\quad ![sender] = @ \circ \langle wr\_resp \rangle,$ send response
$\quad ![r] = Tail(@)]$ remove from channel
ELSE $\land\ channel' = [channel$ EXCEPT $![r] = Tail(@)]$ no response should be sent back
$\land$ UNCHANGED $data$

A secondary replica processes a write request from primary. The secondary fully executes it and sends a reply back to sender. We don't simulate dropping of the request by *sec.* because there's no need to, given the general *TransLoss*.

$SecondaryWrite(r,\ o,\ w,\ ver,\ sender) \triangleq$ changes $\langle data,\ channel,\ stat,\ resps \rangle$
$\land\ \lor\ SecondaryContinuesWrite(r,\ o,\ w,\ ver,\ sender)$

No need to drop the write at the secondaries. This drop will happen due to *TransLoss*().
$\lor\ DropWrite(r,\ o\ w)\ \backslash *$ OR: drop the request alltogether
$\land$ UNCHANGED $\langle stat,\ resps \rangle$

$A(n\ ex -$)primary processes a response to a write request from replica sender. The response might already be too late (it has already timed out), or it might be or might be in-time.

$ProcessWriteResp(r,\ o,\ w,\ sender,\ ack) \triangleq$ changes $\langle data,\ channel,\ stat,\ resps \rangle$
$\land\ resps' = [resps$ EXCEPT $![r][o][sender] =$
IF $@ =$ "timeout" $\lor\ @ =$ "n/a" THEN $@$ resp came too late, it's already expired
ELSE $ack]$ either *ok* or some system-specific error

$\land\ channel' = [channel$ EXCEPT $![r] = Tail(@)]$
$\land$ UNCHANGED $\langle data,\ stat \rangle$

A replica processes a message in its incoming channel. Depending on the type of message (update req, update resp), replica acts according to the three functions above.

$\_ProcessMessage(r) \triangleq$
$\land\ channel[r] \neq \langle\rangle$ I have a message to process

$\land$ LET $m \triangleq Head(channel[r])$IN
$\land$ IF $m.type =$ "wrReq" THEN
IF $m.sender = Client$ THEN upate request from client to a prim
Enters here only when *Client* writes are not a 0-stage action
$PrimaryWrite\ (r,\ m.object,\ m.w,\ m.version)$
TODO: *NB*: This function sends a *Reply*, while the others don't – BUG

ELSE It's a *wrReq* from a primary to a secondary
$SecondaryWrite(r,\ m.object,\ m.w,\ m.version,\ m.sender)$

ELSE IF $m.type =$ "wrResp" THEN
$ProcessWriteResp(r,\ m.object,\ m.w,\ m.sender,\ m.ack)$
ELSE
$\land\ Print($"BUGGGG!!!! Wrong message type!", $m) \neq \langle\rangle$

$\land$ UNCHANGED $\langle master,\ cache \rangle$

$ProcessMessage \triangleq$
$\exists\, r \in Rep :$
$\wedge\ stat[r].phase = \text{“alive”}$ r must be alive to process a message
$\wedge\ \_ProcessMessage(r)$

A replica (either a current primary or not) finishes a write that it started some time ago (when it considered itself to be a primary).

$\_FinishWrite(r,\ o) \triangleq$
LET
$NotFinishedWrReqs(m) \triangleq m.object \neq o \vee m.sender \neq r$
this message does not refer to a write on object o
initiated by replica r. So, I don't declare it finished
$NotFinishedWrResps(m) \triangleq m.object \neq o$
IN
$\wedge\ stat[r].phase = \text{“alive”}$ allow writes to finish even for those started by dead replicas
TODO: *Shouldn*'t do this for Blue??????
$\wedge\ stat[r].lock[o] = \text{“busy”}$ r has indeed a write started
$\wedge\ \forall\, r1 \in Rep : resps[r][o][r1] \neq \text{“waiting”}$ all the replicas have submitted
their responses, or have timed-out
or weren't supposed to answer
$\wedge\ ReplyFinishedWrite(r,\ o,\ stat[r].in\_progress[o])$
should send the appropriate *Reply*, based on state of the *resps*
but not change any of master, data, cache, etc.!
Re-initialize all write-related state, to prepare the replica for the next write
$\wedge\ stat' = [stat$ EXCEPT $![r].lock[o] = \text{“rdy”},$ r is ready to accept new updates
(if it's still prim, of course)
$![r].in\_progress[o] = [val \mapsto NoWrite,\ version \mapsto 0]]$
re-init the responses, to prepare r for the next write
$\wedge\ resps' = [resps$ EXCEPT $![r][o] = [r1 \in Rep \mapsto \text{“waiting”}]]$
Remove ALL messages related to this object. Leave the rest intact
$\wedge\ channel' = [r1 \in Rep \mapsto$ IF $r1 \neq r$ THEN $SelectSeq(channel[r1],\ NotFinishedWrReqs)$
ELSE $SelectSeq(channel[r],\ NotFinishedWrResps)]$
I think that the above is not actually cheating, b/c you can imagine
an implementation where replicas would simply identify older
updates (via a sequence number, which is anyway necessary for
keeping the order of updates.)
$\wedge$ UNCHANGED $\langle master,\ data,\ cache\rangle$

$FinishWrite \triangleq$
$\exists\, r \in Rep,\ o \in Object : \_FinishWrite(r,\ o)$

$ReplicaActions \triangleq$
$\vee\ ReplicaDeath$

$\lor$ *ReadVersion*
$\lor$ *ProcessMessage*
$\lor$ *FinishWrite* a replica finishes a write it has started

---

CLIENT actions.

Client wants to perform write w on object o. The request is performed in a 0-stage fashion. The method models caching of replica locations in the client – BUT IT'S NOT CORRECT !!! *TODO TODO*.

$\_CliWrite(prim,\ o,\ w) \triangleq$
  $\land$ $PrimaryWrite(prim,\ o,\ w,\ master.objects[o].version)$
  $\land$ UNCHANGED $\langle master,\ cache \rangle$
$CliWrite(o,\ w) \triangleq$
  $\exists\, r \in Rep :$
    $\land$ $stat[r].phase =$ "alive"
    $\land$ $\_CliWrite(r,\ o,\ w)$

In the most general case, a read is typically performed on the local copy of some replica that stores that object (is either a *sec* or a primary). How you choose that replica depends on the protocol. In *GFS*, it's any replica. In Blue, it's gotta be prm The semantic of the read is that I read a whole object.

$\_CliRead(r,\ o) \triangleq$
  $\land$ LET $val \triangleq StoreRead(r,\ o)$
    IN $Reply(r,$ "rd", $val)$
  $\land$ UNCHANGED $\langle master,\ cache,\ data,\ stat,\ resps,\ channel \rangle$

$CliRead(o) \triangleq$
  $\exists\, r \in Rep :$
    $\land$ $stat[r].phase =$ "alive"
    $\land$ $master.health[r] =$ "alive" ASSUMPTION : *Uncomment* this if you want to test read-last-successful-X ($X =$ write/append). This ensures I don't read from a live but stale replica.
    $\land$ $r \in cache[r][o].sec \cup \{cache[r][o].prim\}$
      r considers itself either a *sec* or primary of o
    $\land$ $\_CliRead(r,\ o)$

$ClientActions \triangleq$
  $\exists\, o \in Object :$
    $\lor \exists\, w \in WriteType : CliWrite(o,\ w)$
    $\lor$ $CliRead(o)$

$Next \triangleq$ $\lor$ *MasterActions*
  $\lor$ *TimeActions*
  $\lor$ *ReplicaActions*
  $\lor$ *ClientActions*

$Spec \triangleq \land Init \land \Box[Next]_{vars}$

Invariants

$AllInvariants \triangleq$
$\quad \land TypeInvariant$

THEOREM $Spec \Rightarrow \Box AllInvariants$