Vanish: Increasing Data Privacy with Self-Destructing Data

Roxana Geambasu
Yoshi Kohno
Amit Levy
Hank Levy

University of Washington

Outline

Part 1: Introducing Self-Destructing Data

Part 2: Vanish Architecture and Implementation

Part 3: Evaluation and Applications

Motivating Problem: Data Lives Forever

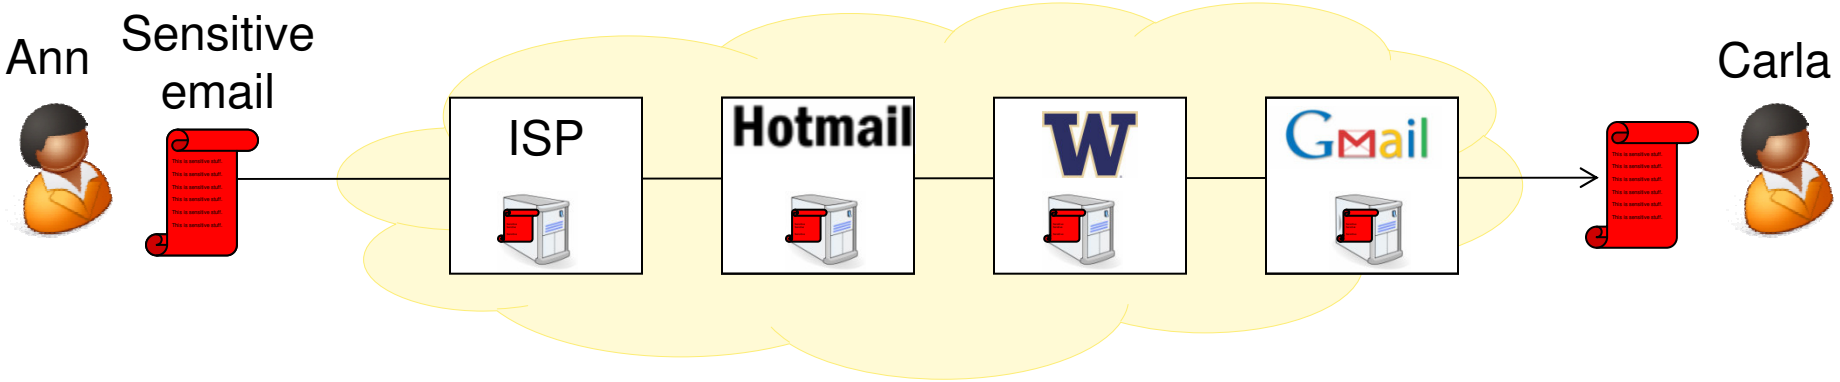

How can Ann delete her sensitive email?

- She doesn't know where all the copies are
- Services may retain data for long after user tries to delete

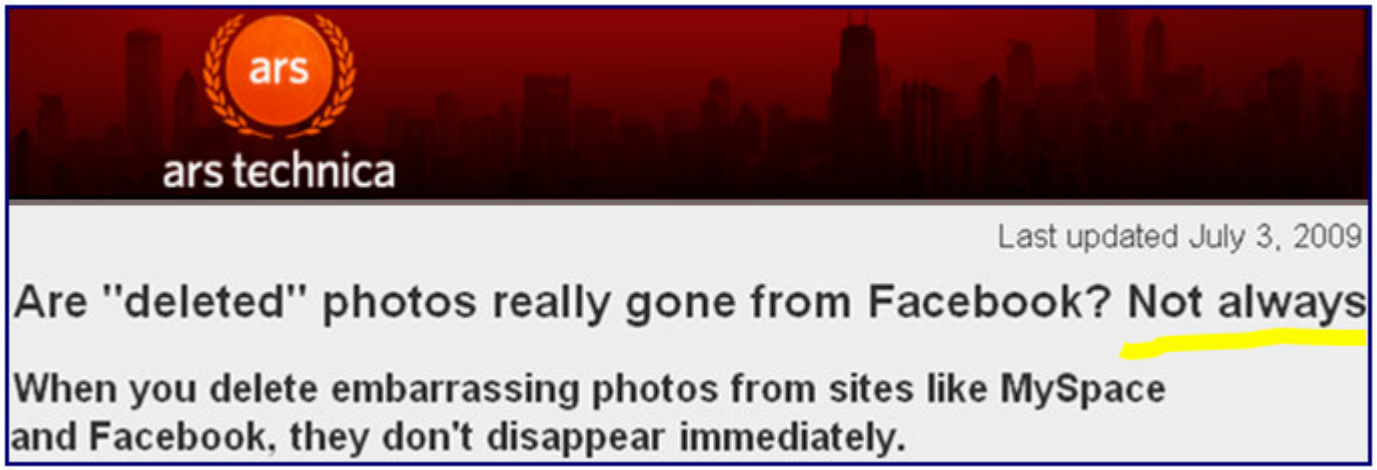

Archived Copies Can Resurface Years Later

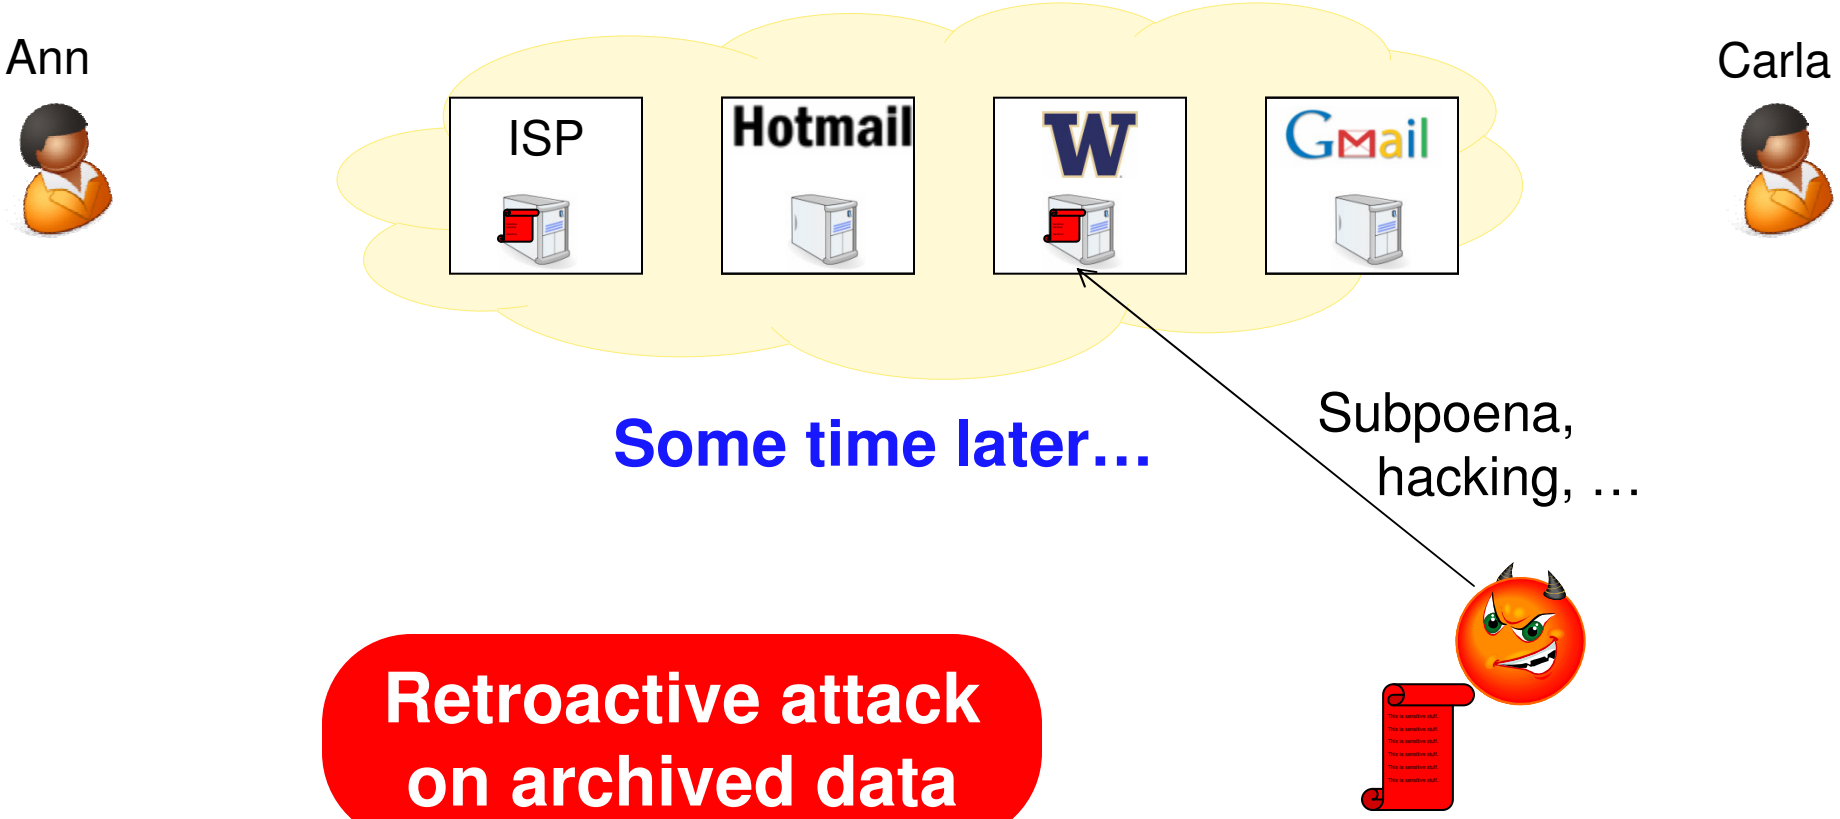

The Retroactive Attack

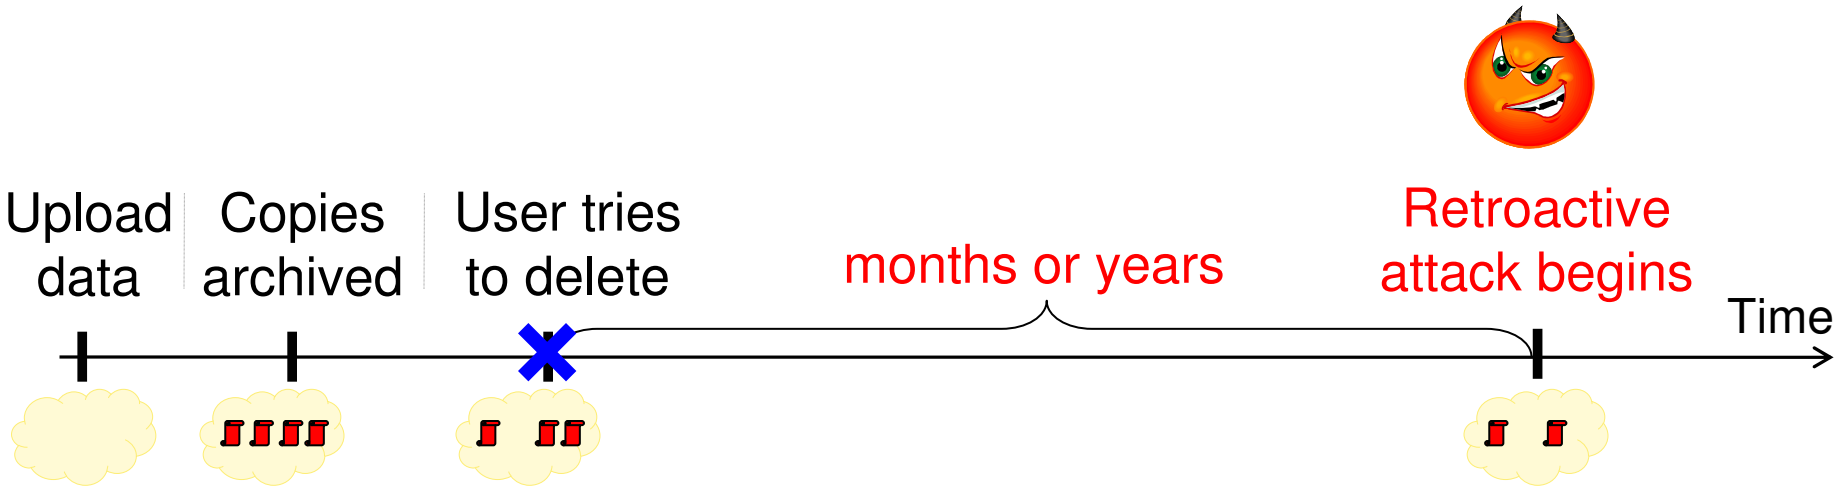

Why Not Use Encryption (e.g., PGP)?

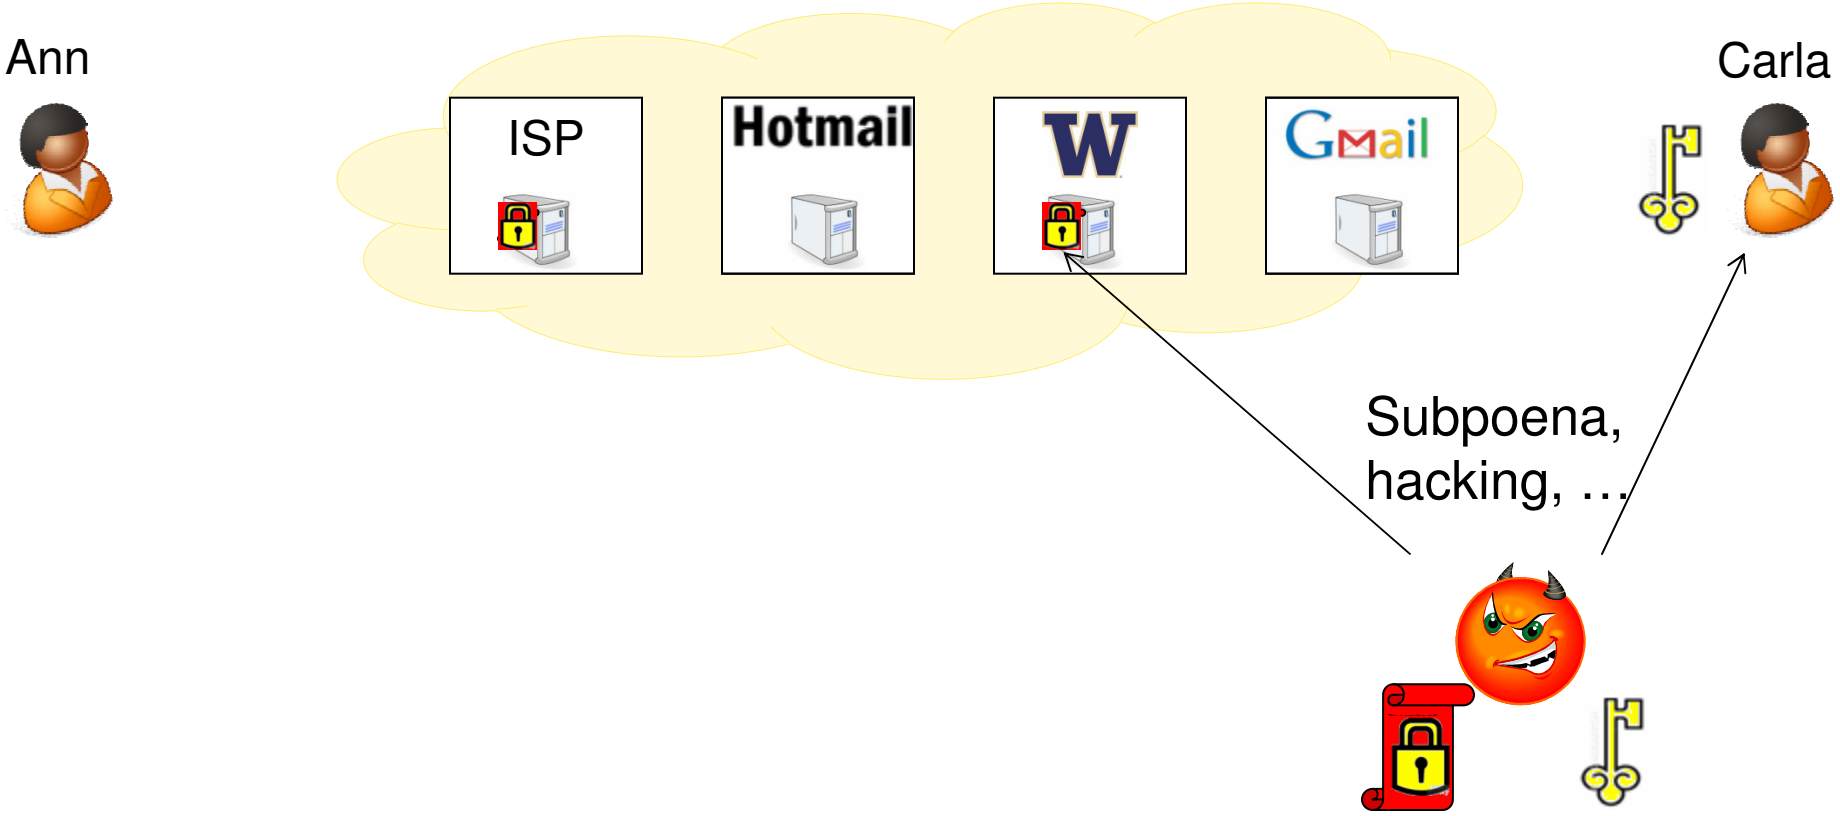

Why Not Use Encryption (e.g., PGP)?

cnet news

February 26, 2009 1:30 PM PST

Judge orders defendant to decrypt PGP-protected laptop

A federal judge has ordered a criminal defendant to decrypt his hard drive by typing i
a ruling

V3.co.uk formerly vnunet.com

UK police can now demand encryption keys

vnunet.com, 03 Oct 2007

People in the UK who encrypt their data are now obliged by law to give up the encryption keys to law enforcement officials if requested under the Regulation of Investigatory Powers Act 2000 (RIP Act).

Why Not Use a Centralized Service?

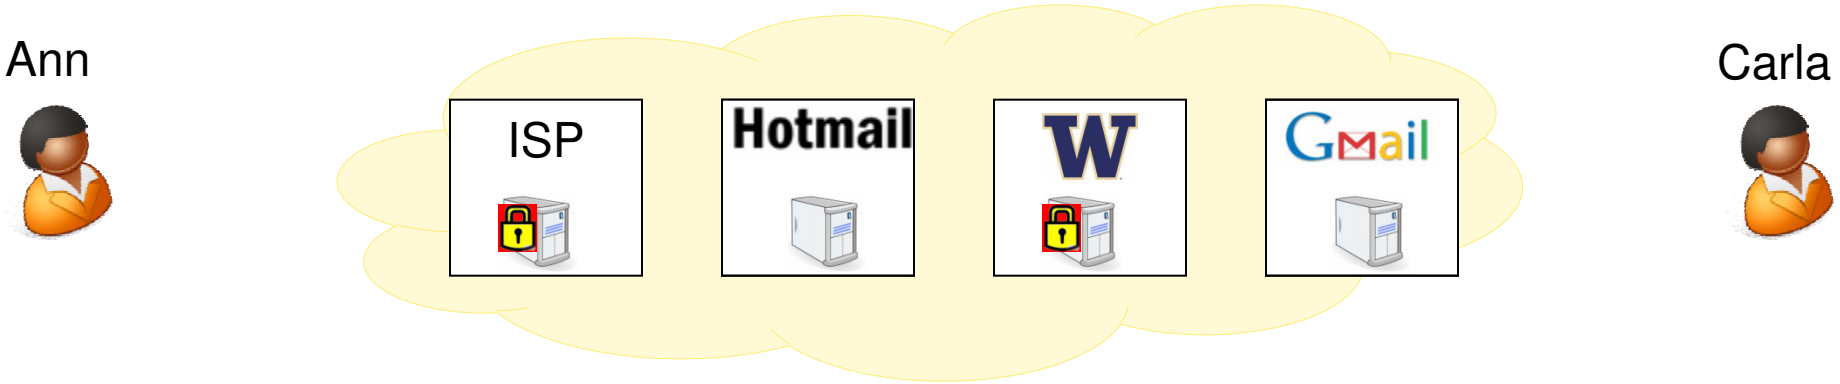

Why Not Use a Centralized Service?

Ann

Carla

WIRED

November 7, 2007 | 3:39 pm

Encrypted E-Mail Company Hushmail Spills to Feds

Hushmail, a longtime provider of encrypted web-based email, markets itself by saying that "not even a Hushmail employee with access to our servers can read your encrypted e-mail, since each message is uniquely encoded before it leaves your computer."

But it turns out that statement seems not to apply to individuals targeted by government agencies that are able to convince a Canadian court to serve a court order on the company.

Centralized Service

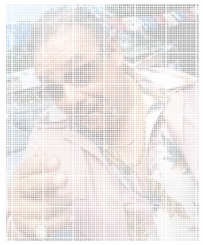

"Trust us: we'll help you delete your data on time."

Backdoor agreement

The Problem: Two Huge Challenges for Privacy

1. Data lives forever
   - On the web: emails, Facebook photos, Google Docs, blogs, …
   - In the home: disks are cheap, so no need to ever delete data
   - In your pocket: phones and USB sticks have GBs of storage
2. Retroactive disclosure of both data and user keys has become commonplace
   - Hackers
   - Misconfigurations
   - Legal actions
   - Border seizing
   - Theft
   - Carelessness


The Washington Post

Palin's Yahoo! Account Hacked

A group of computer hackers said yesterday they accessed a Yahoo! e-mail account of Alaska Gov. Sarah Palin, the Republican vice presidential nominee, publishing some of her private communications [...]

The Problem: Two Huge Challenges for Privacy

1. Data lives forever
   - On the web: emails, Facebook photos, Google Docs, blogs, …
   - In the home: disks are cheap, so no need to ever delete data
   - In your pocket: phones and USB sticks have GBs of storage
2. Retroactive disclosure of both data and user keys has become commonplace
   - Hackers
   - Misconfigurations
   - Legal actions
   - Border seizing
   - Theft
   - Carelessness

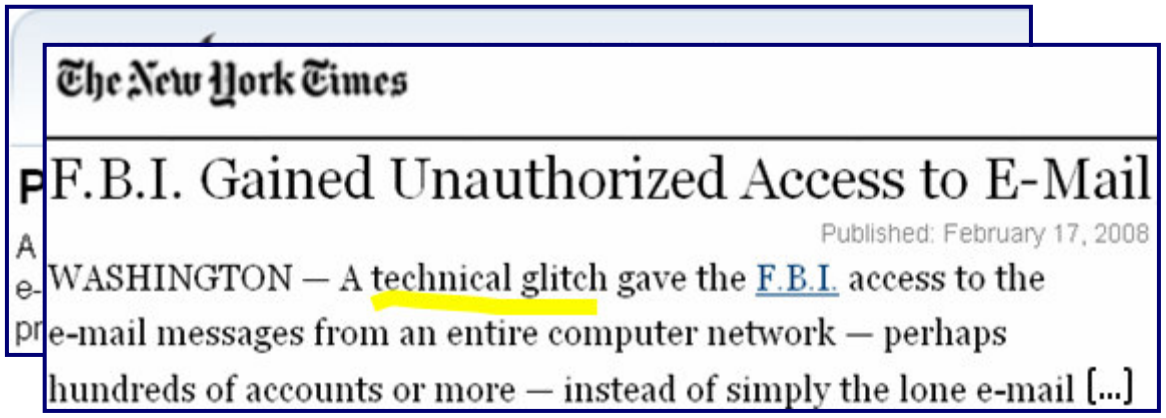

The New York Times

F.B.I. Gained Unauthorized Access to E-Mail

Published: February 17, 2008

WASHINGTON — A technical glitch gave the F.B.I. access to the e-mail messages from an entire computer network — perhaps hundreds of accounts or more — instead of simply the lone e-mail [...]

The Problem: Two Huge Challenges for Privacy

1. Data lives forever
   - On the web: emails, Facebook photos, Google Docs, blogs, …
   - In the home: disks are cheap, so no need to ever delete data
   - In your pocket: phones and USB sticks have GBs of storage
2. Retroactive disclosure of both data and user keys has become commonplace
   - Hackers
   - Misconfigurations
   - Legal actions
   - Border seizing
   - Theft
   - Carelessness

The Problem: Two Huge Challenges for Privacy

1. Data lives forever
   - On the web: emails, Facebook photos, Google Docs, blogs, …
   - In the home: disks are cheap, so no need to ever delete data
   - In your pocket: phones and USB sticks have GBs of storage
2. Retroactive disclosure of both data and user keys has become commonplace
   - Hackers
   - Misconfigurations
   - Legal actions
   - Border seizing
   - Theft
   - Carelessness

Question:

Can we empower users with control of data lifetime?

Answer:

Self-destructing data

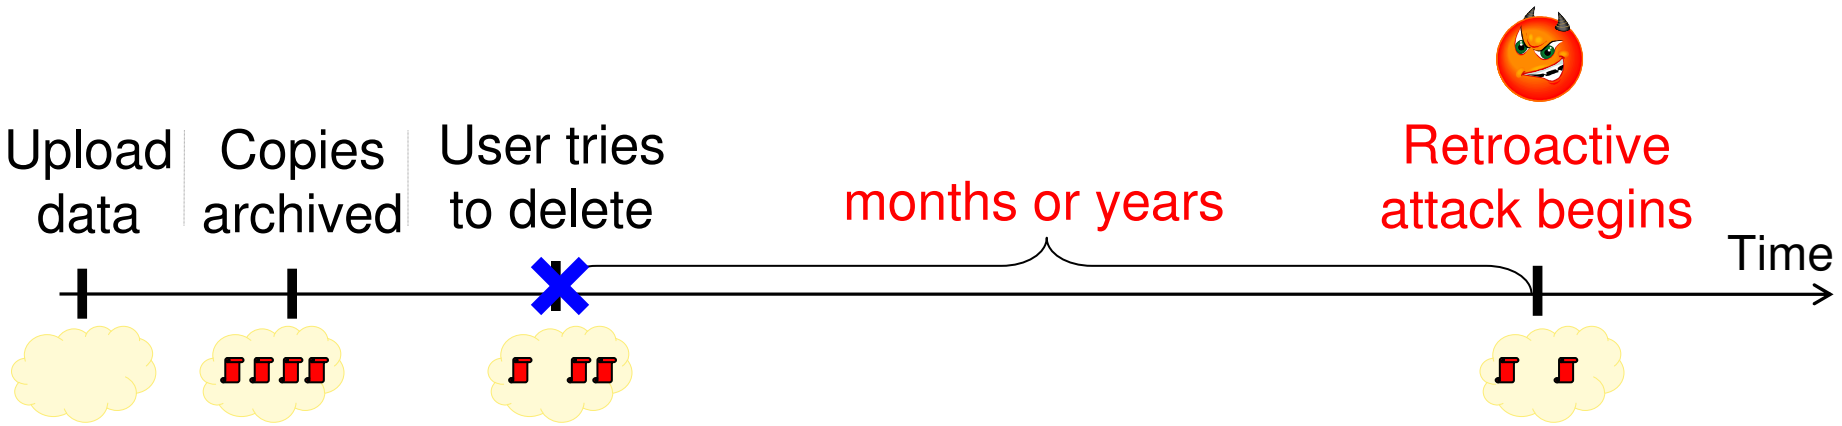

Question:

Can we empower users with control of data lifetime?

Answer:

Self-destructing data

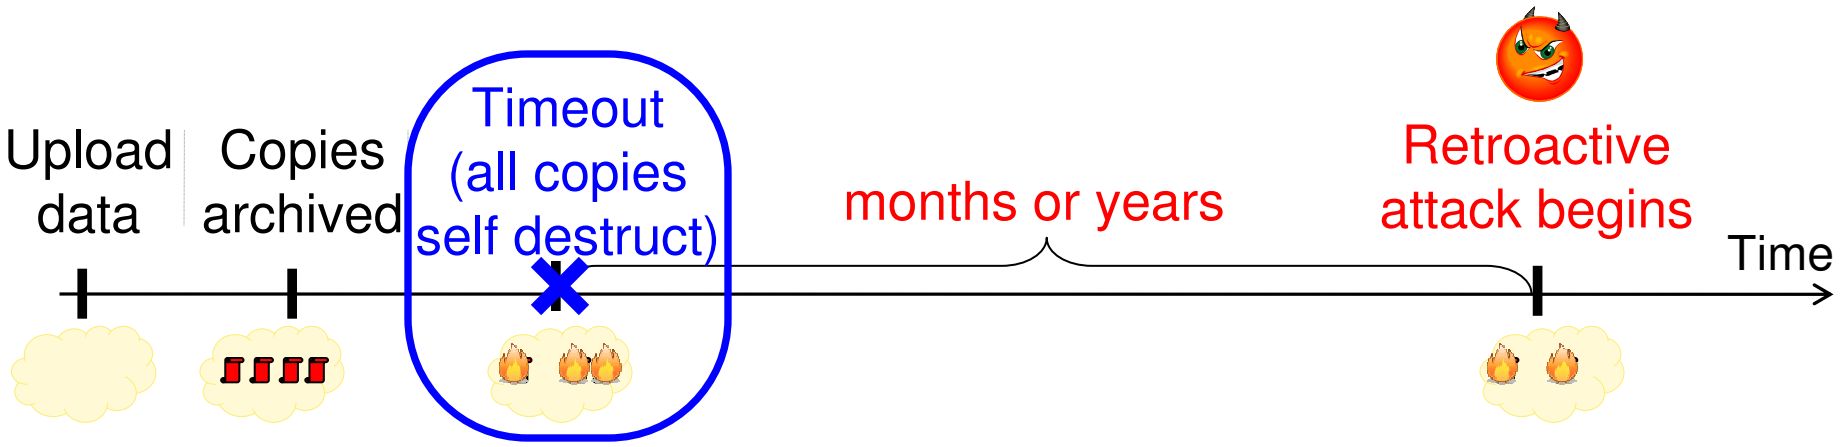

Self-Destructing Data Model

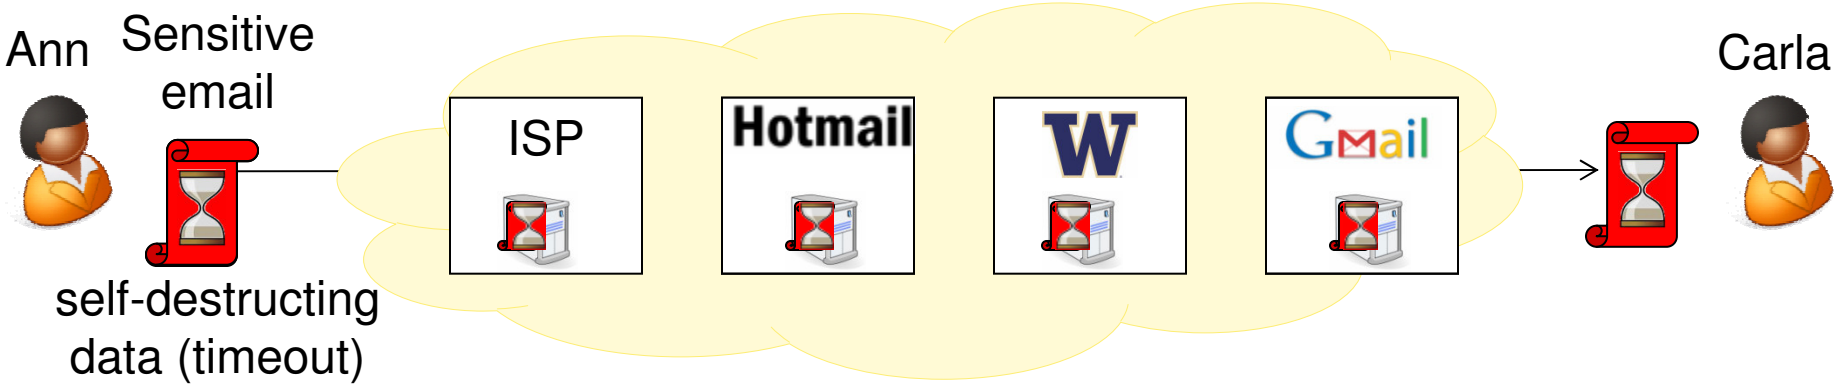

1. Until timeout, users can read original message

Self-Destructing Data Model

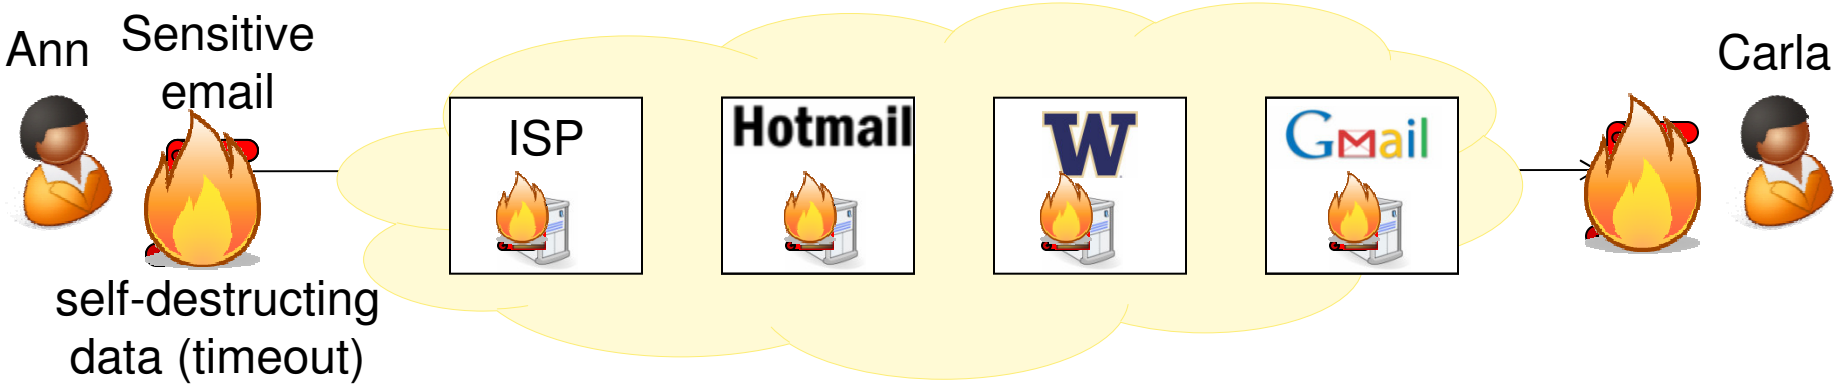

1. Until timeout, users can read original message
2. After timeout, all copies become permanently unreadable
   - 2.1. even for attackers who obtain an archived copy & user keys
   - 2.2. without requiring explicit delete action by user/services
   - 2.3. without having to trust any centralized services

Self-Destructing Data Model

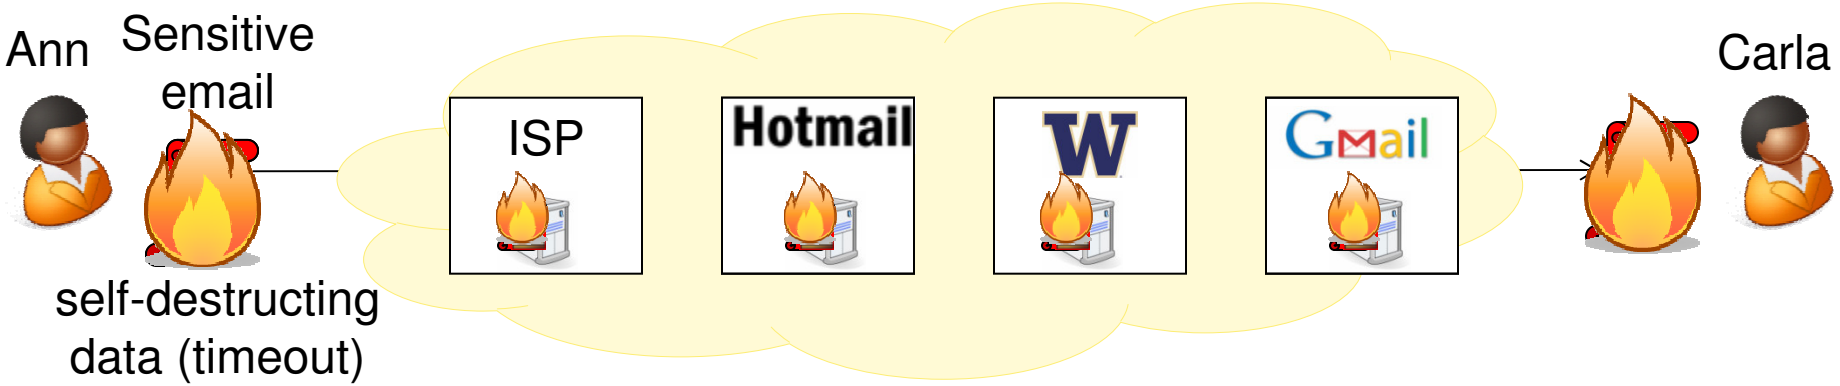

Goals of Self-Destructing Data

1. Until timeout, users can read original message
2. After timeout, all copies become permanently unreadable
   - 2.1. even for attackers who obtain an archived copy & user keys
   - 2.2. without requiring explicit delete action by user/services
   - 2.3. without having to trust any centralized services

Outline

Part 1: Introducing Self-Destructing Data

Part 2: Vanish Architecture and Implementation

Part 3: Evaluation and Applications

Vanish: Self-Destructing Data System

- Traditional solutions are not sufficient for self-destructing data goals:
  - PGP
  - Centralized data management services
  - Forward-secure encryption
  - …
- Let's try something completely new!

Idea:
Leverage P2P systems

P2P 101: Intro to Peer-To-Peer Systems

- A system composed of individually-owned computers that make a portion of their resources available directly to their peers without intermediary managed hosts or servers. [~wikipedia]

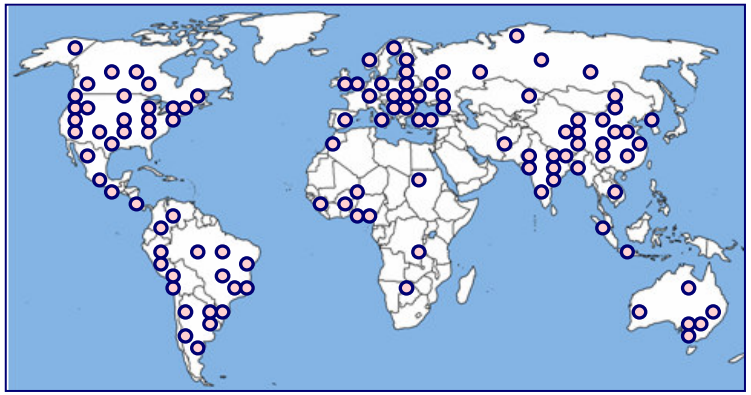

Important P2P properties (for Vanish):

- Huge scale – millions of nodes
- Geographic distribution – hundreds of countries
- Decentralization – individually-owned, no single point of trust
- Constant evolution – nodes constantly join and leave

Distributed Hashtables (DHTs)

- Hashtable data structure implemented on a P2P network
  - Get and put (index, value) pairs
  - Each node stores part of the index space

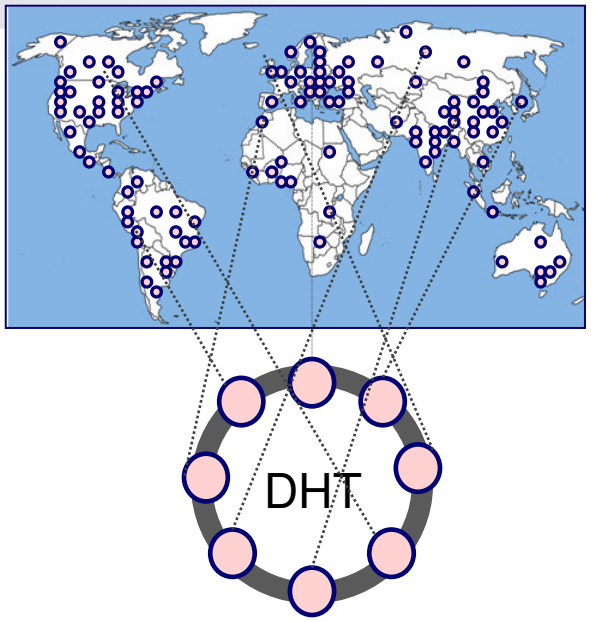

Logical structure

- DHTs are part of many file sharing systems:
  - Vuze, Mainline, KAD
  - Vuze has ~1.5M simultaneous nodes in ~190 countries

- Vanish leverages DHTs to provide self-destructing data
  - One of few applications of DHTs outside of file sharing

How Vanish Works: Data Encapsulation

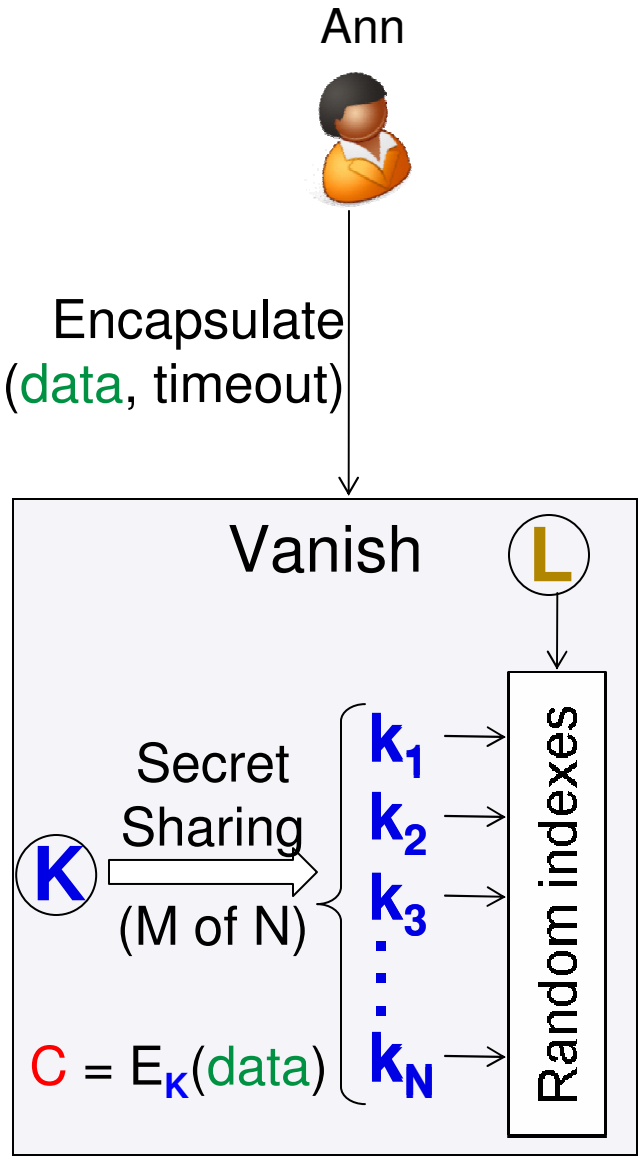

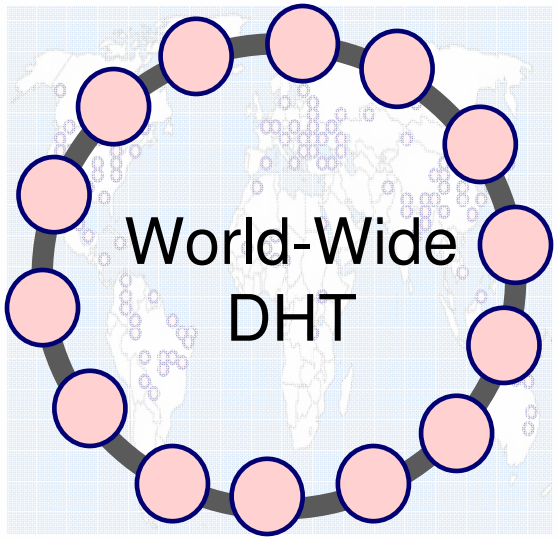

How Vanish Works: Data Encapsulation

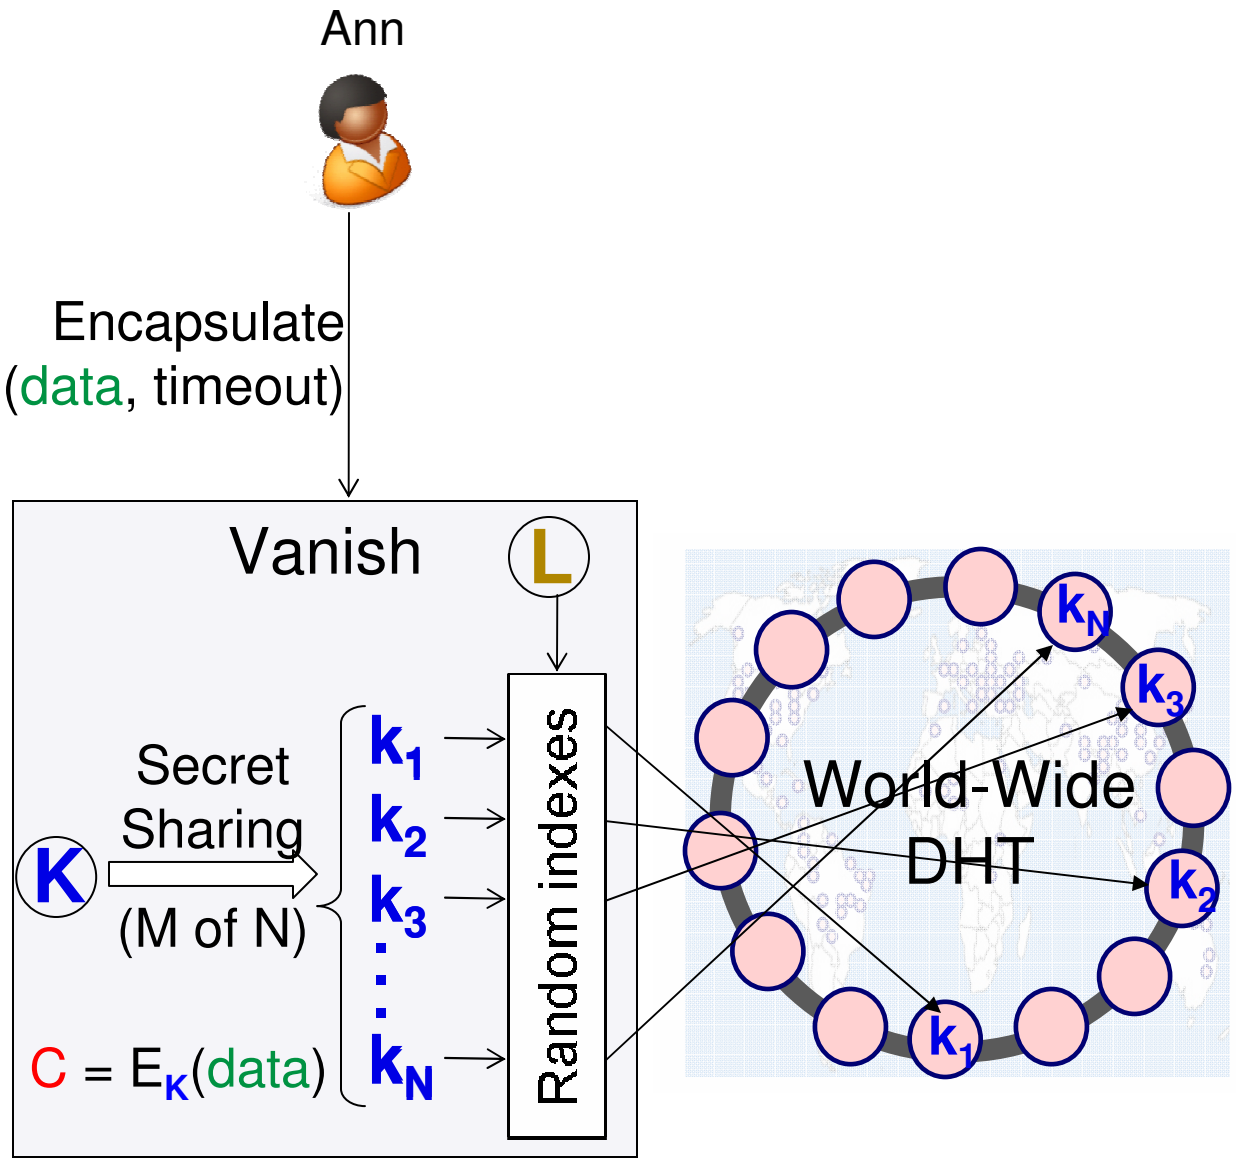

How Vanish Works: Data Encapsulation

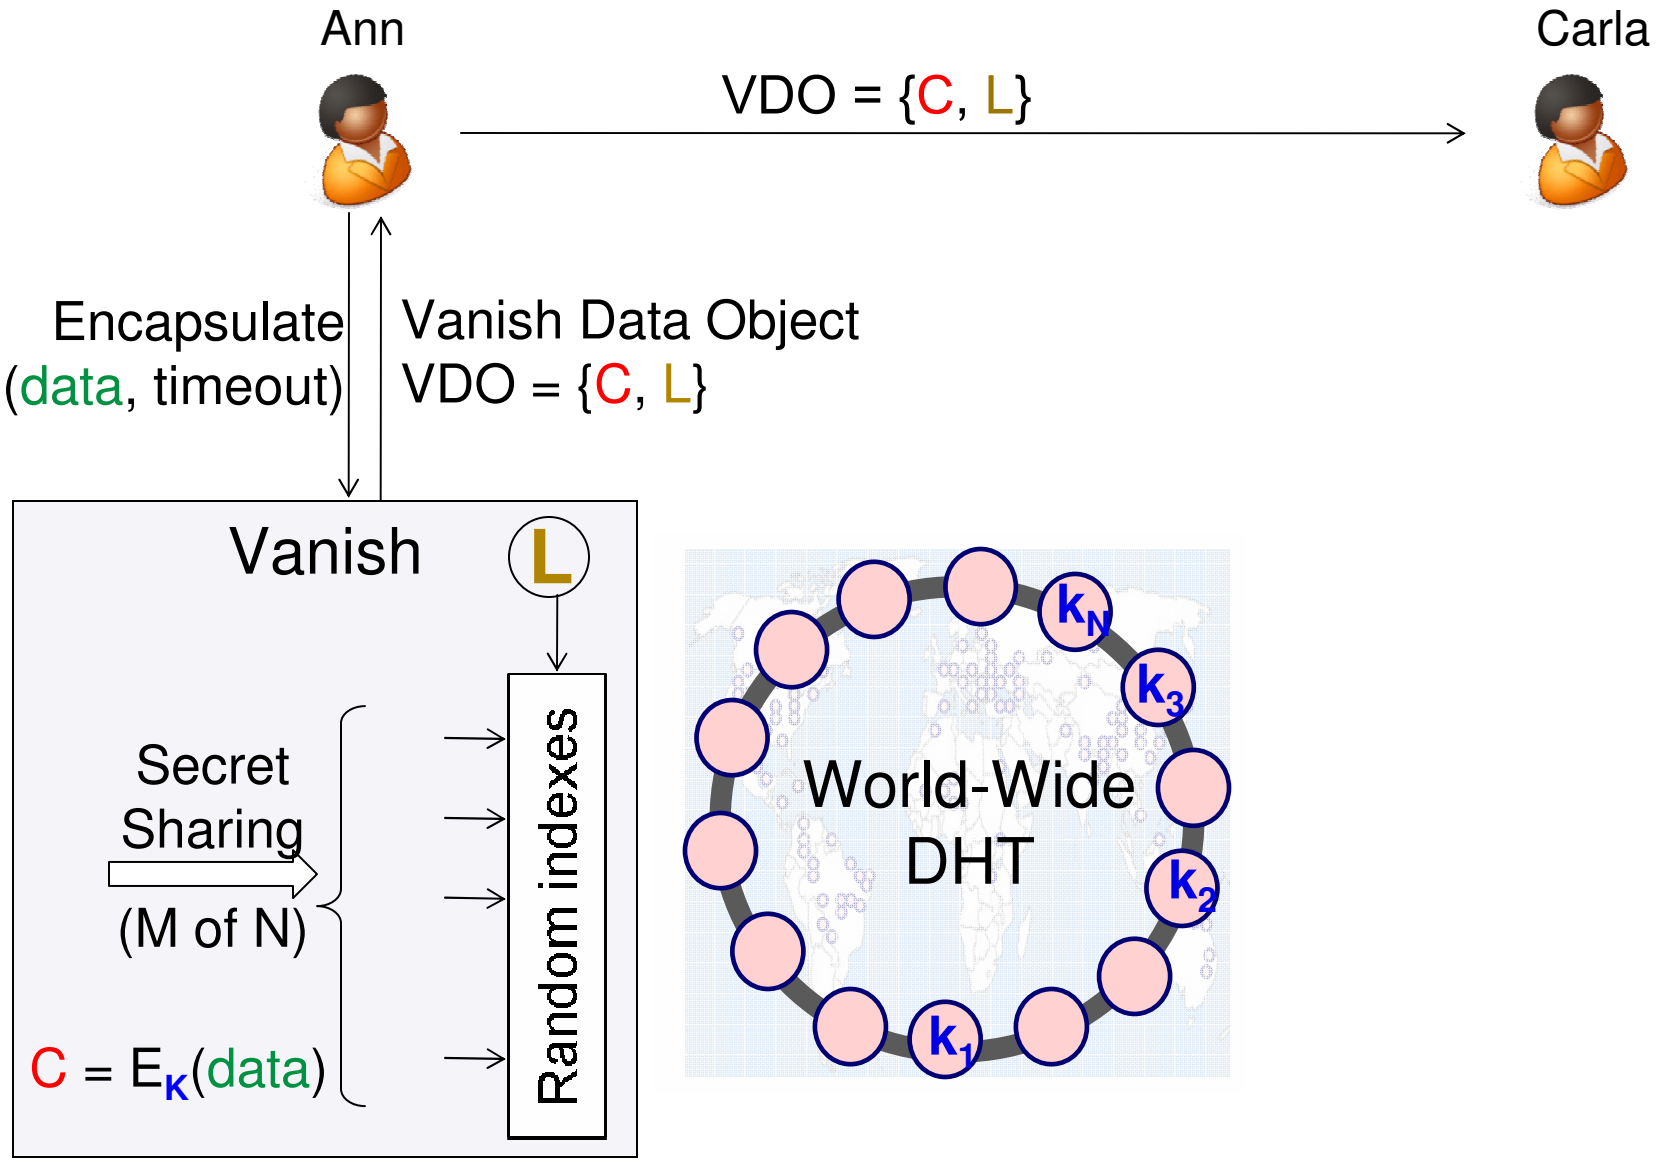

How Vanish Works: Data Decapsulation

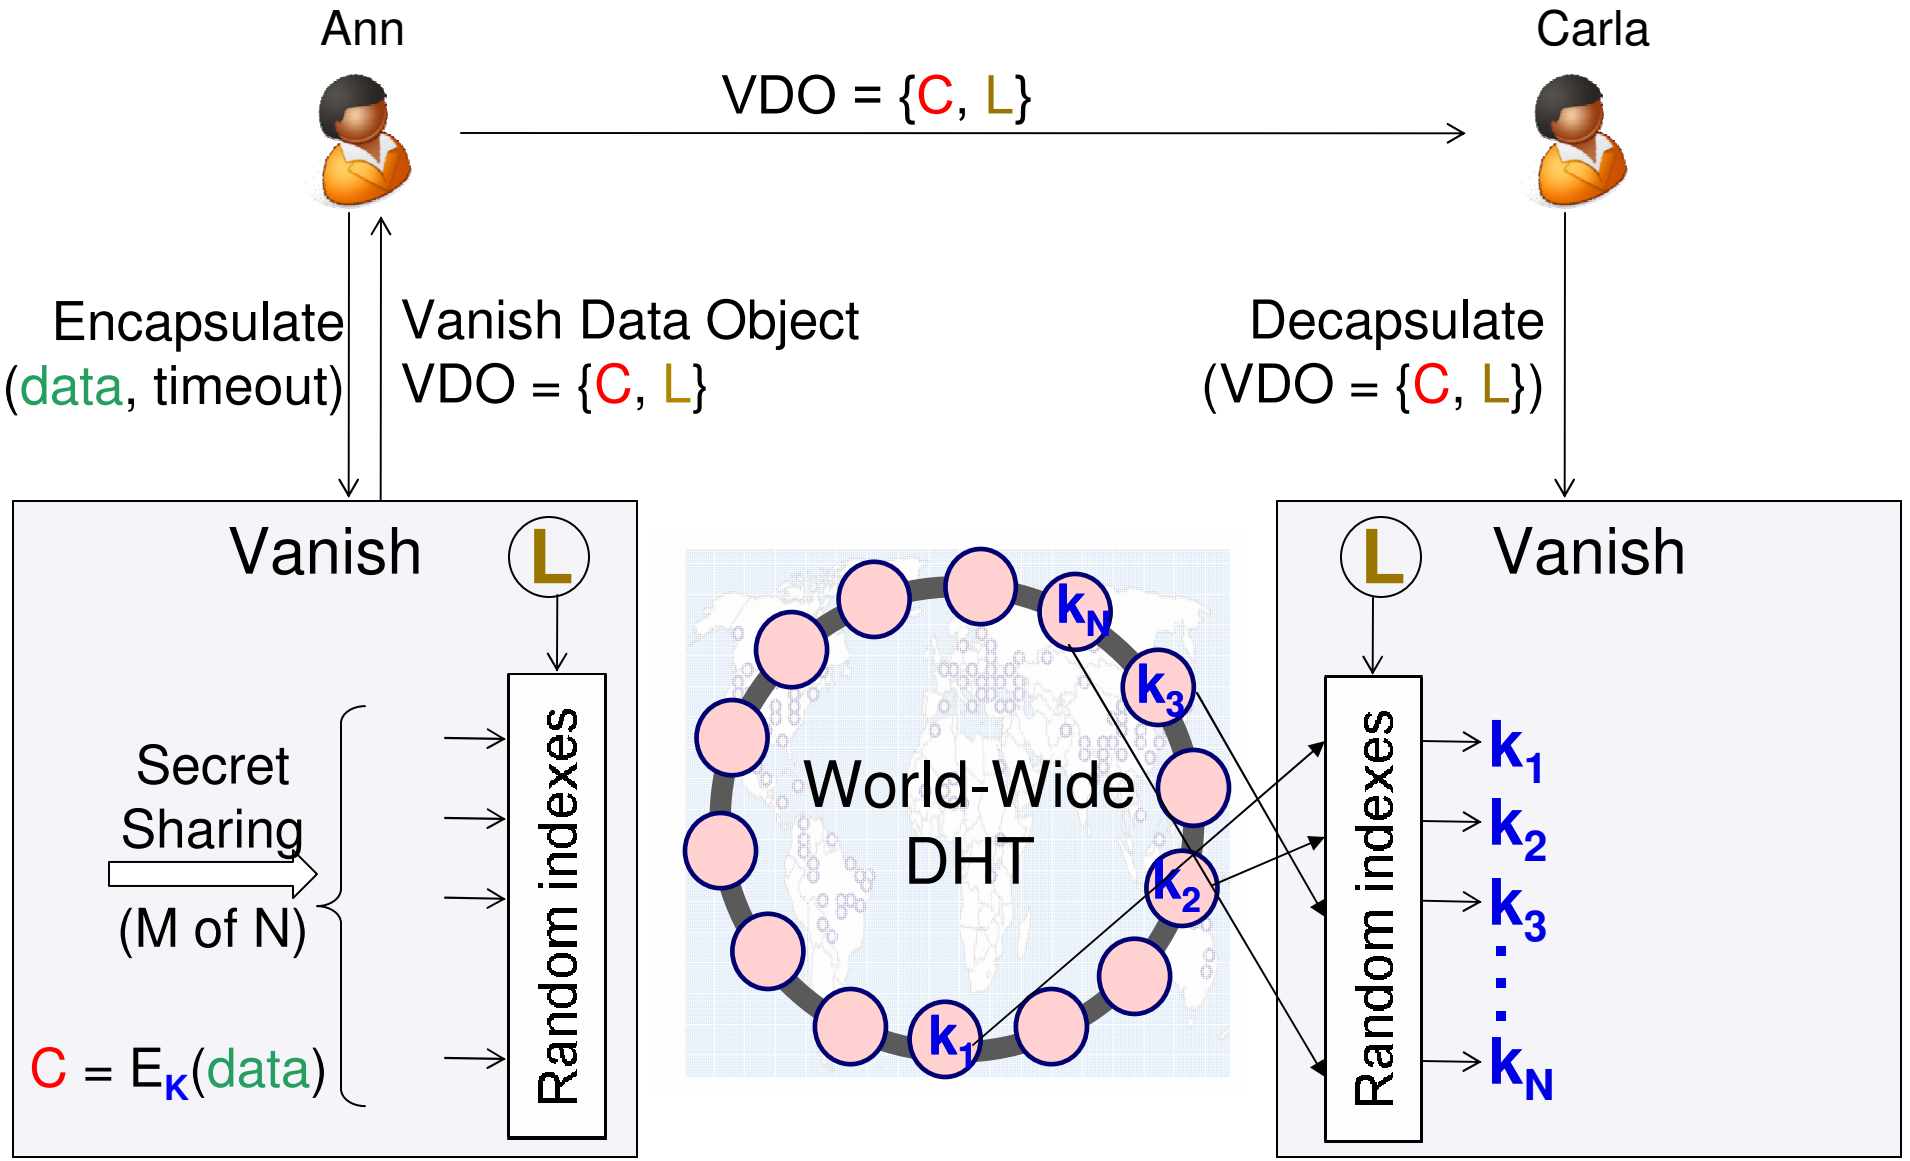

How Vanish Works: Data Decapsulation

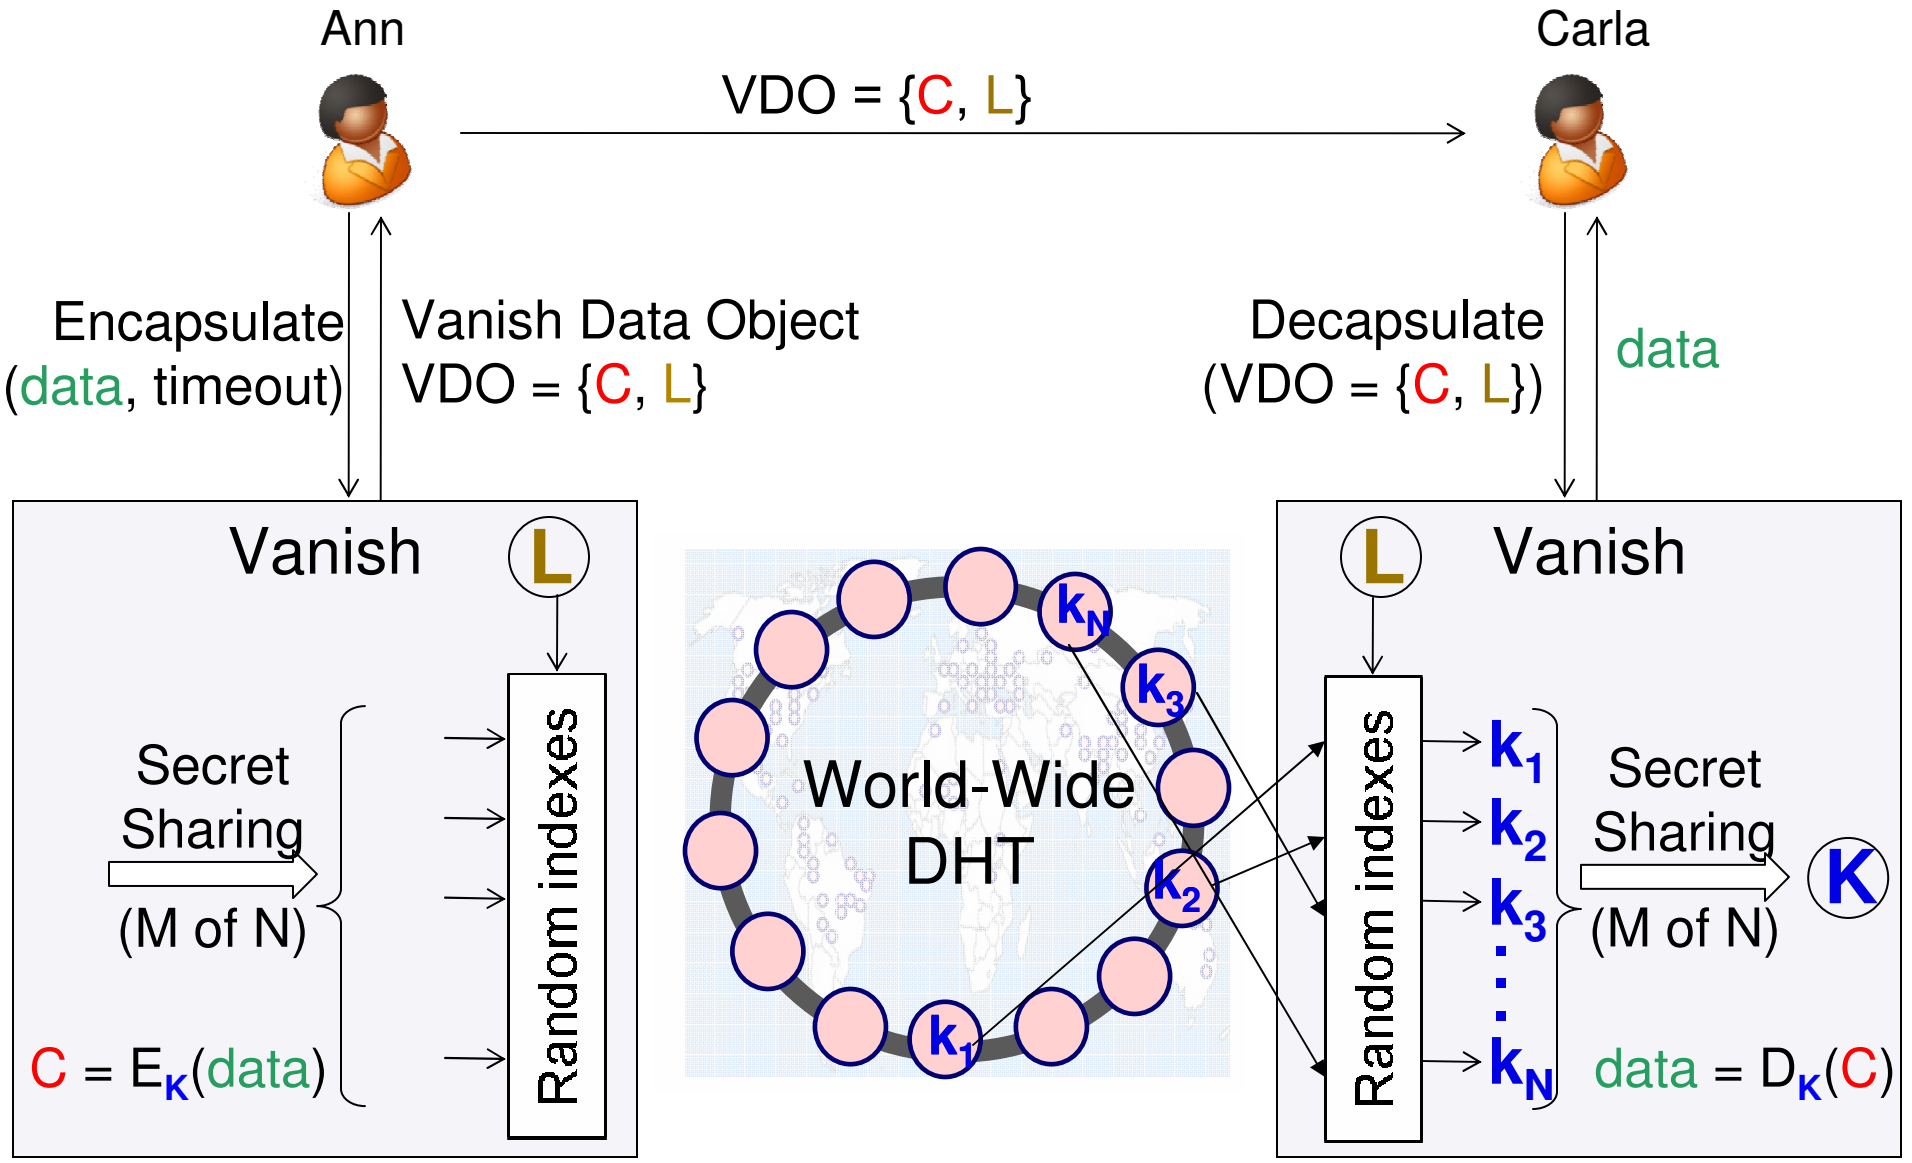

How Vanish Works: Data Decapsulation

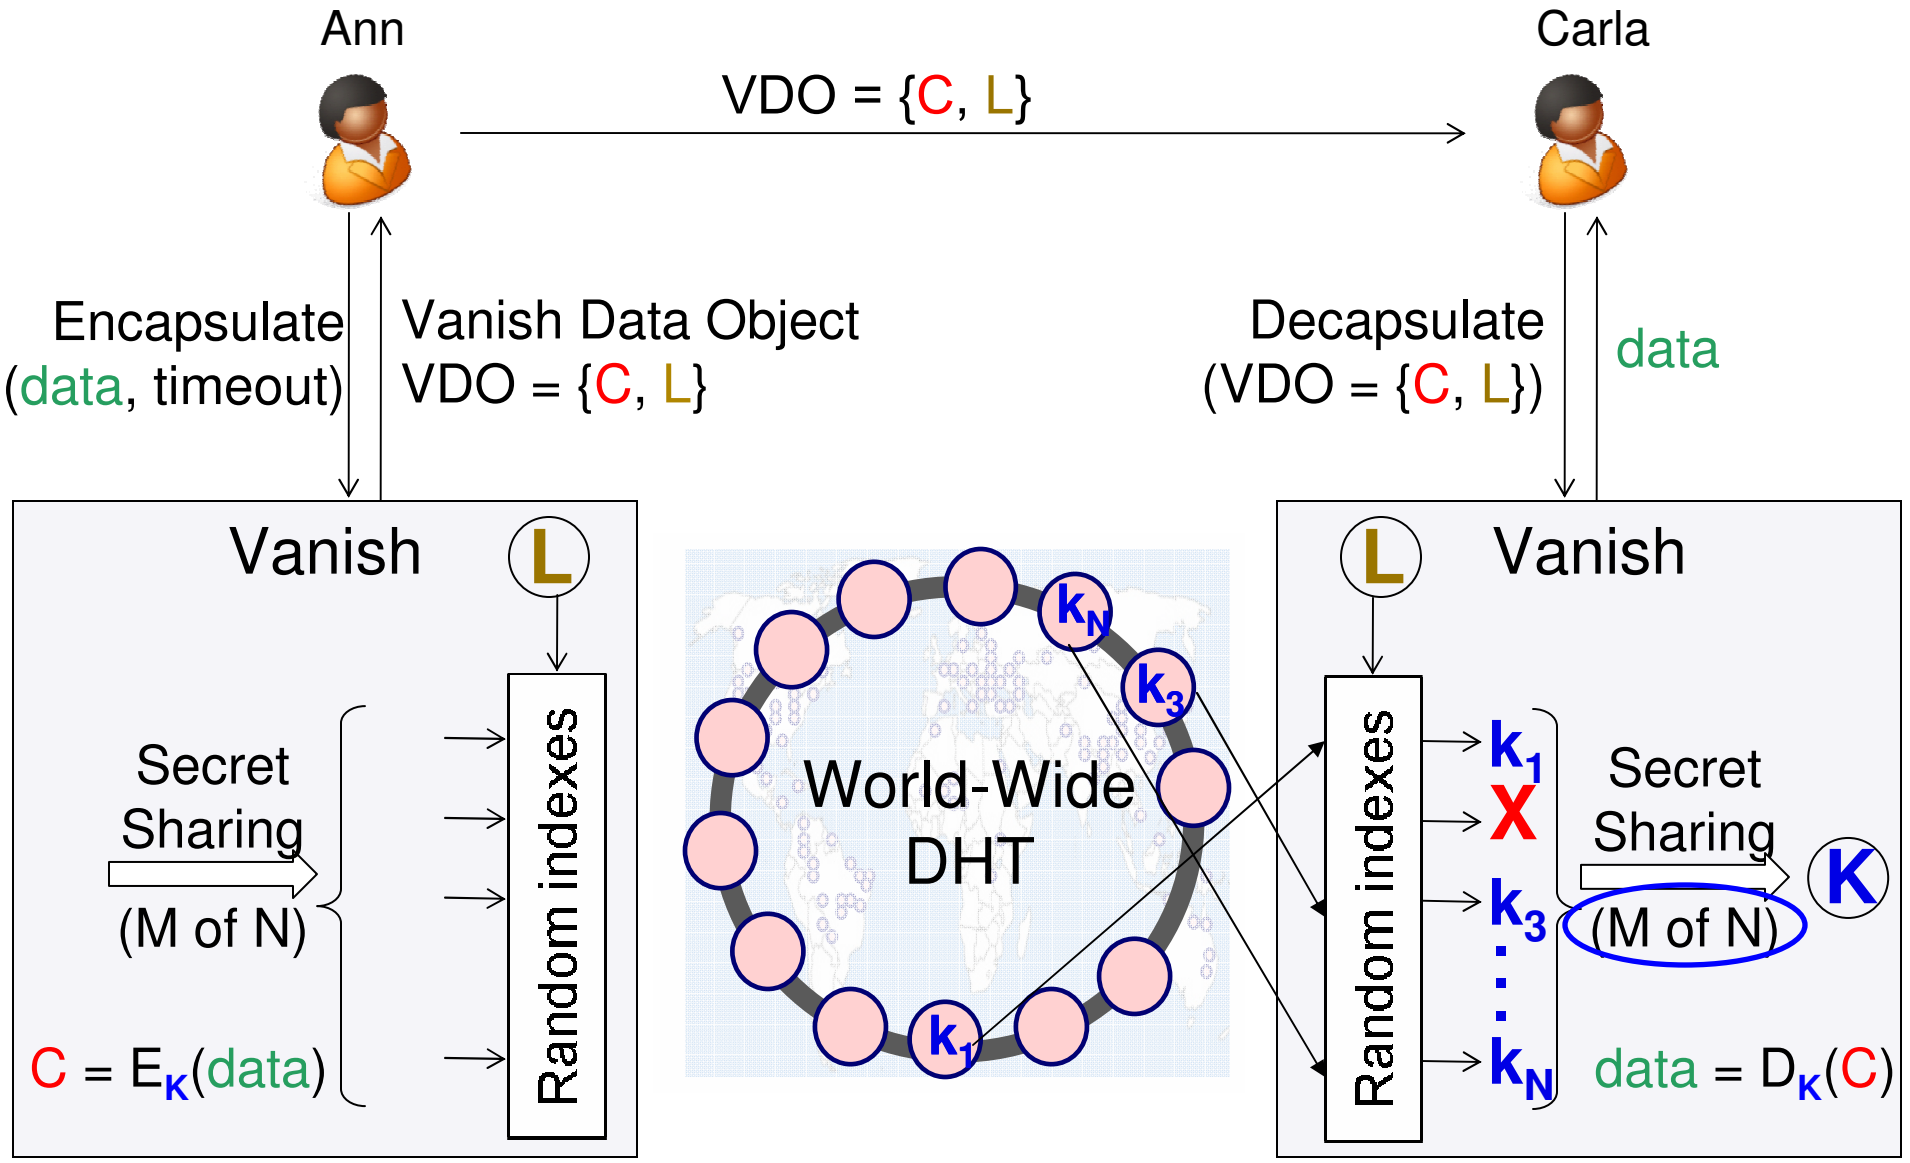

How Vanish Works: Data Timeout

- The DHT loses key pieces over time
  - Natural churn: nodes crash or leave the DHT
  - Built-in timeout: DHT nodes purge data periodically

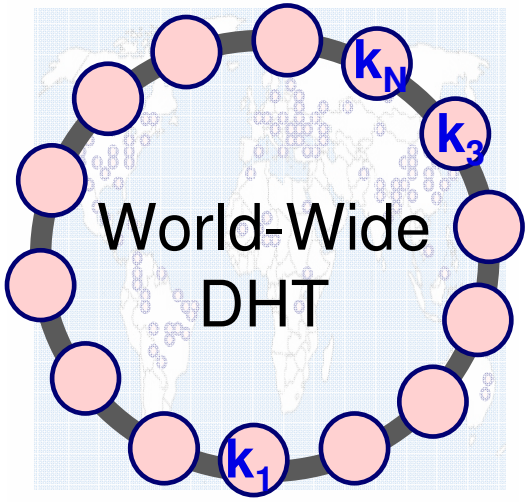

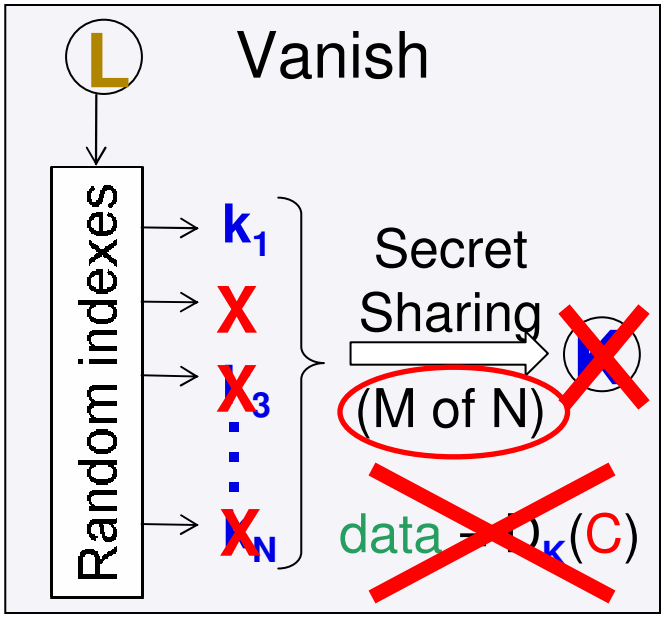

- Key loss makes all data copies permanently unreadable

Outline

Part 1: Introducing Self-Destructing Data

Part 2: Vanish Architecture and Implementation

Part 3: Evaluation and Applications

Evaluation

- Experiments to understand and improve:
  1. data availability before timeout
  2. data unavailability after timeout
  3. performance
  4. security

  In the paper

  Discussed next

- Highest-level results:
  - Secret sharing parameters (N and M) affect availability, timeout, performance, and security
  - Tradeoffs are necessary

Threat Model

- Goal: protect against retroactive attacks on old copies
  - Attackers don't know their target until after timeout
  - Attackers may do non-targeted "pre-computations" at any time

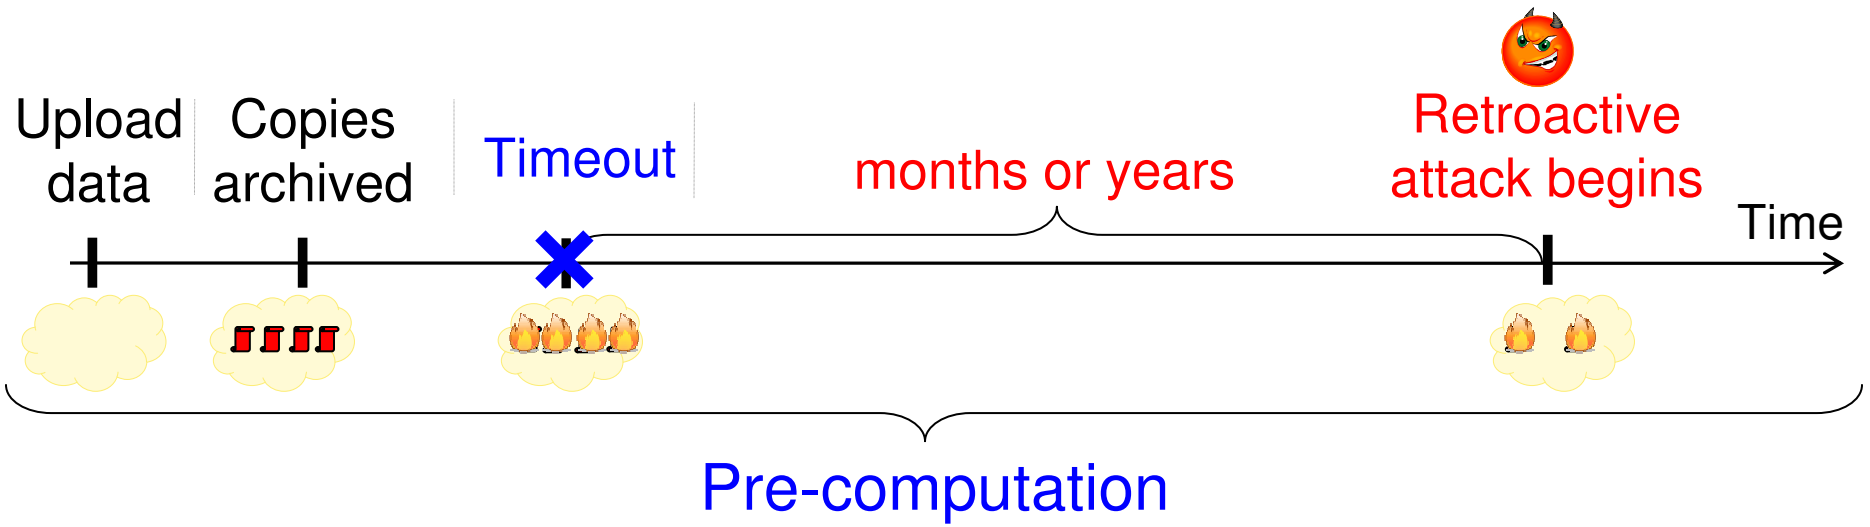

- Communicating parties trust each other
  - E.g., Ann trusts Carla not to keep a plain-text copy

Attack Analysis

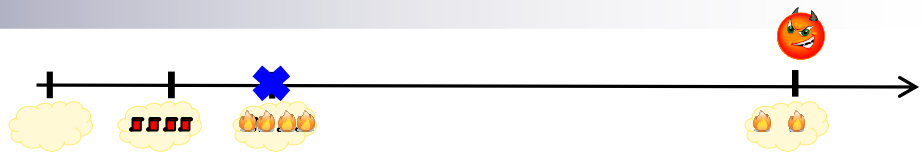

Retroactive Attack	Defense
Obtain data by legal means (e.g., subpoenas)	P2P properties: constant evolution, geographic distribution, decentralization
Gmail decapsulates all VDO emails	Compose with traditional encryption (e.g., PGP)
ISP sniffs traffic	Anonymity systems (e.g., Tor)
DHT eclipse, routing attack	Defenses in DHT literature (e.g., constraints on routing table)
DHT Sybil attack	Defenses in DHT literature; Vuze offers some basic protection
Intercept DHT “get” requests & save results	Vanish obfuscates key share lookups
Capture key pieces from the DHT (pre-computation)	P2P property: huge scale
More (see paper)	

Attack Analysis

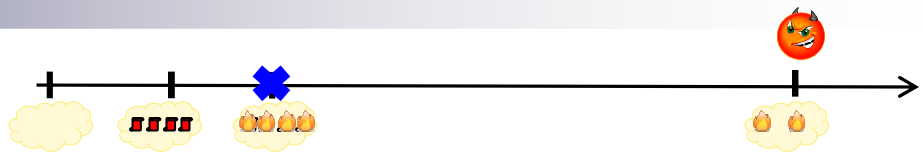

Retroactive Attack	Defense
Obtain data by legal means (e.g., subpoenas)	P2P properties: constant evolution, geographic distribution, decentralization
Gmail decapsulates all VDO emails	Compose with traditional encryption (e.g., PGP)
ISP sniffs traffic	g., Tor)
DHT eclipse, routing	ture (e.g., constraints
DHT Sybil attack	ture; Vuze offers
Intercept DHT "get" requests & save results	Vanish obfuscates key share lookups
Capture key pieces from the DHT (pre-computation)	P2P property: huge scale
More (see paper)	

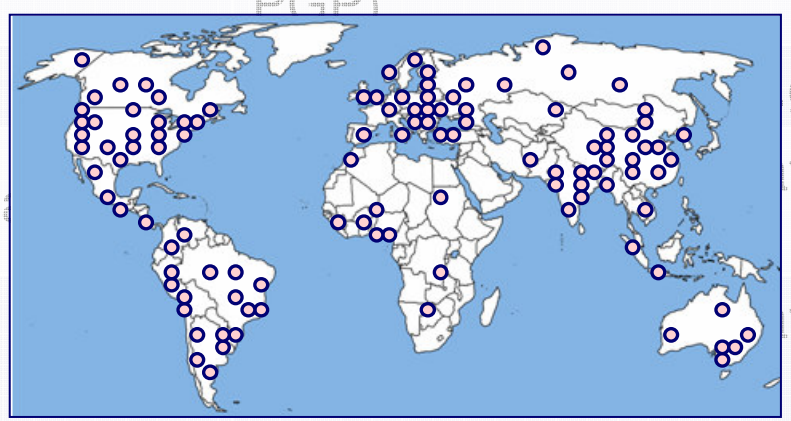

Retroactive Attacks

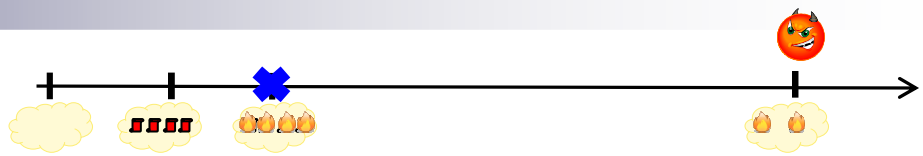

Attack	Defense
Capture any key pieces from the DHT (pre-computation)	P2P property: huge scale

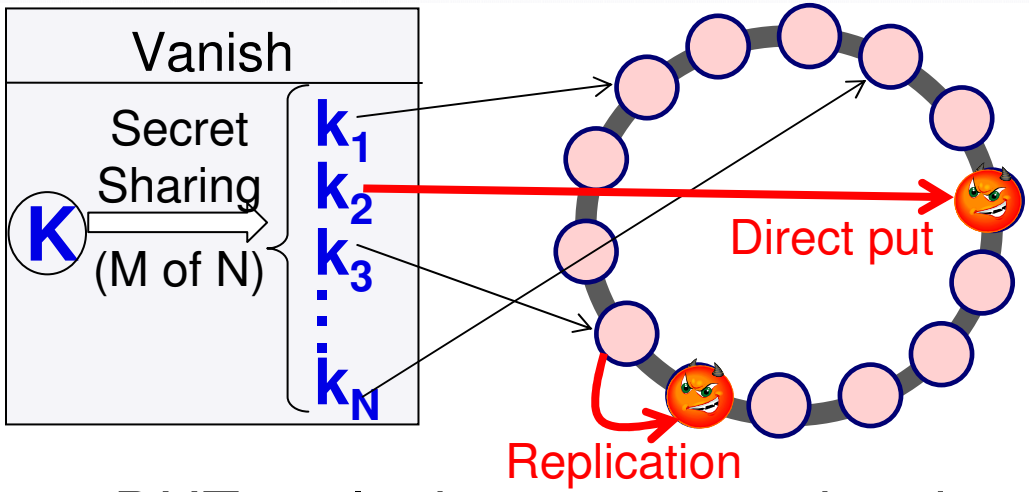

- Given the huge DHT scale, how many nodes does the attacker need to be effective?
- Current estimate:
  - Attacker must join with ~8% of DHT size, for 25% capture
  - There may be other attacks (and defenses)

Vanish Applications

- Self-destructing data & Vanish support many applications

Example applications:

- Firefox plugin
  - Included in our release of Vanish
- Thunderbird plugin
  - Developed by the community two weeks after release ☺
- Self-destructing files
- Self-destructing trash-bin
- …

Firefox Plugin For Vanishing Web Data

- Encapsulate text in any text area in self-destructing VDOs

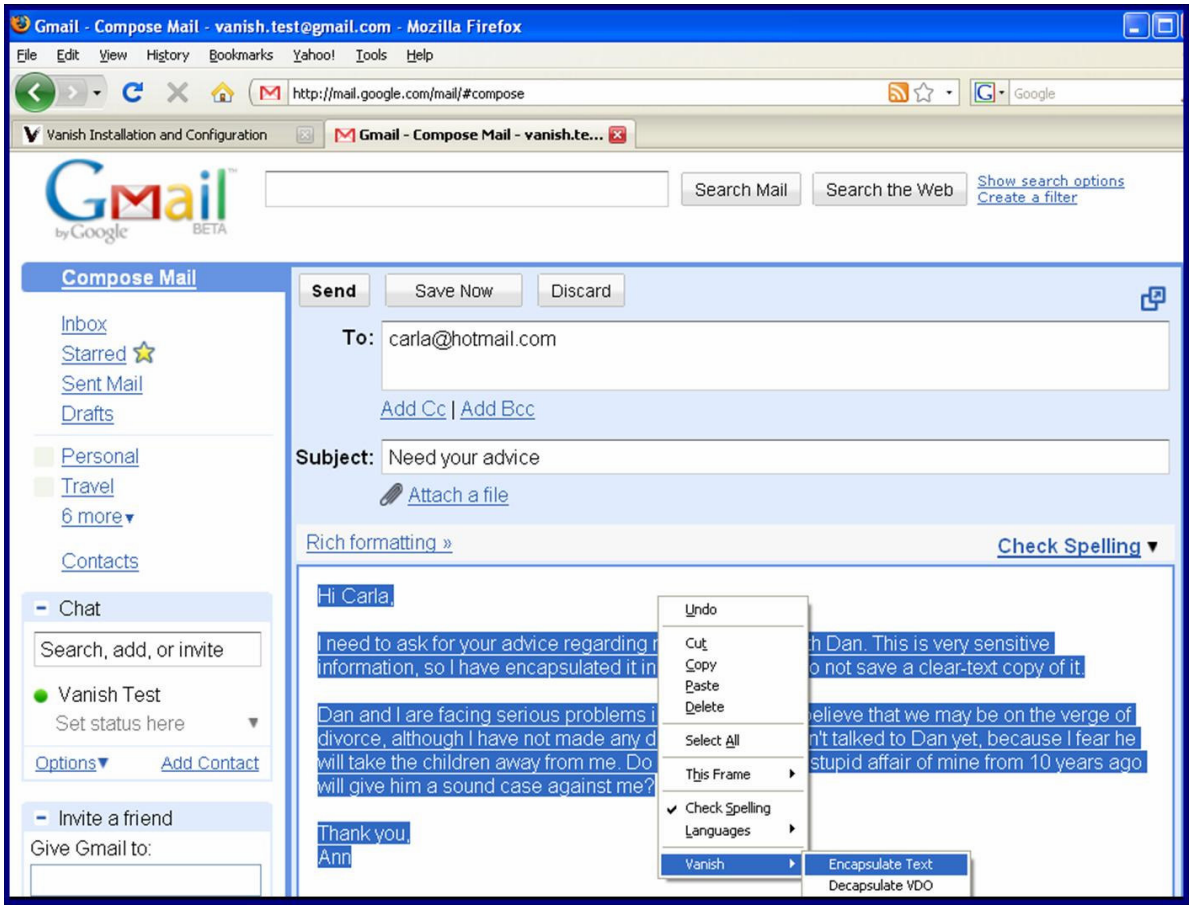

Firefox Plugin For Vanishing Web Data

- Encapsulate text in any text area in self-destructing VDOs

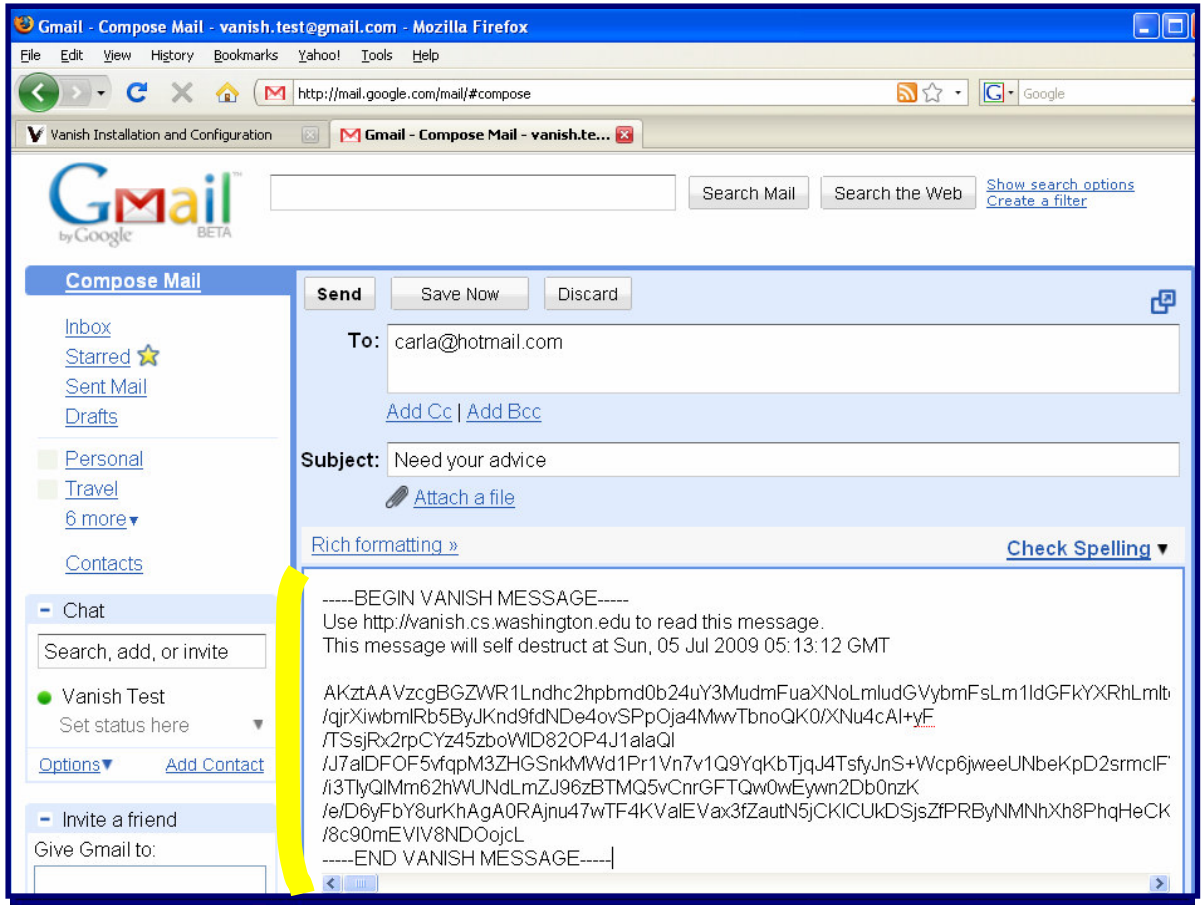

Firefox Plugin For Vanishing Web Data

- Encapsulate text in any text area in self-destructing VDOs

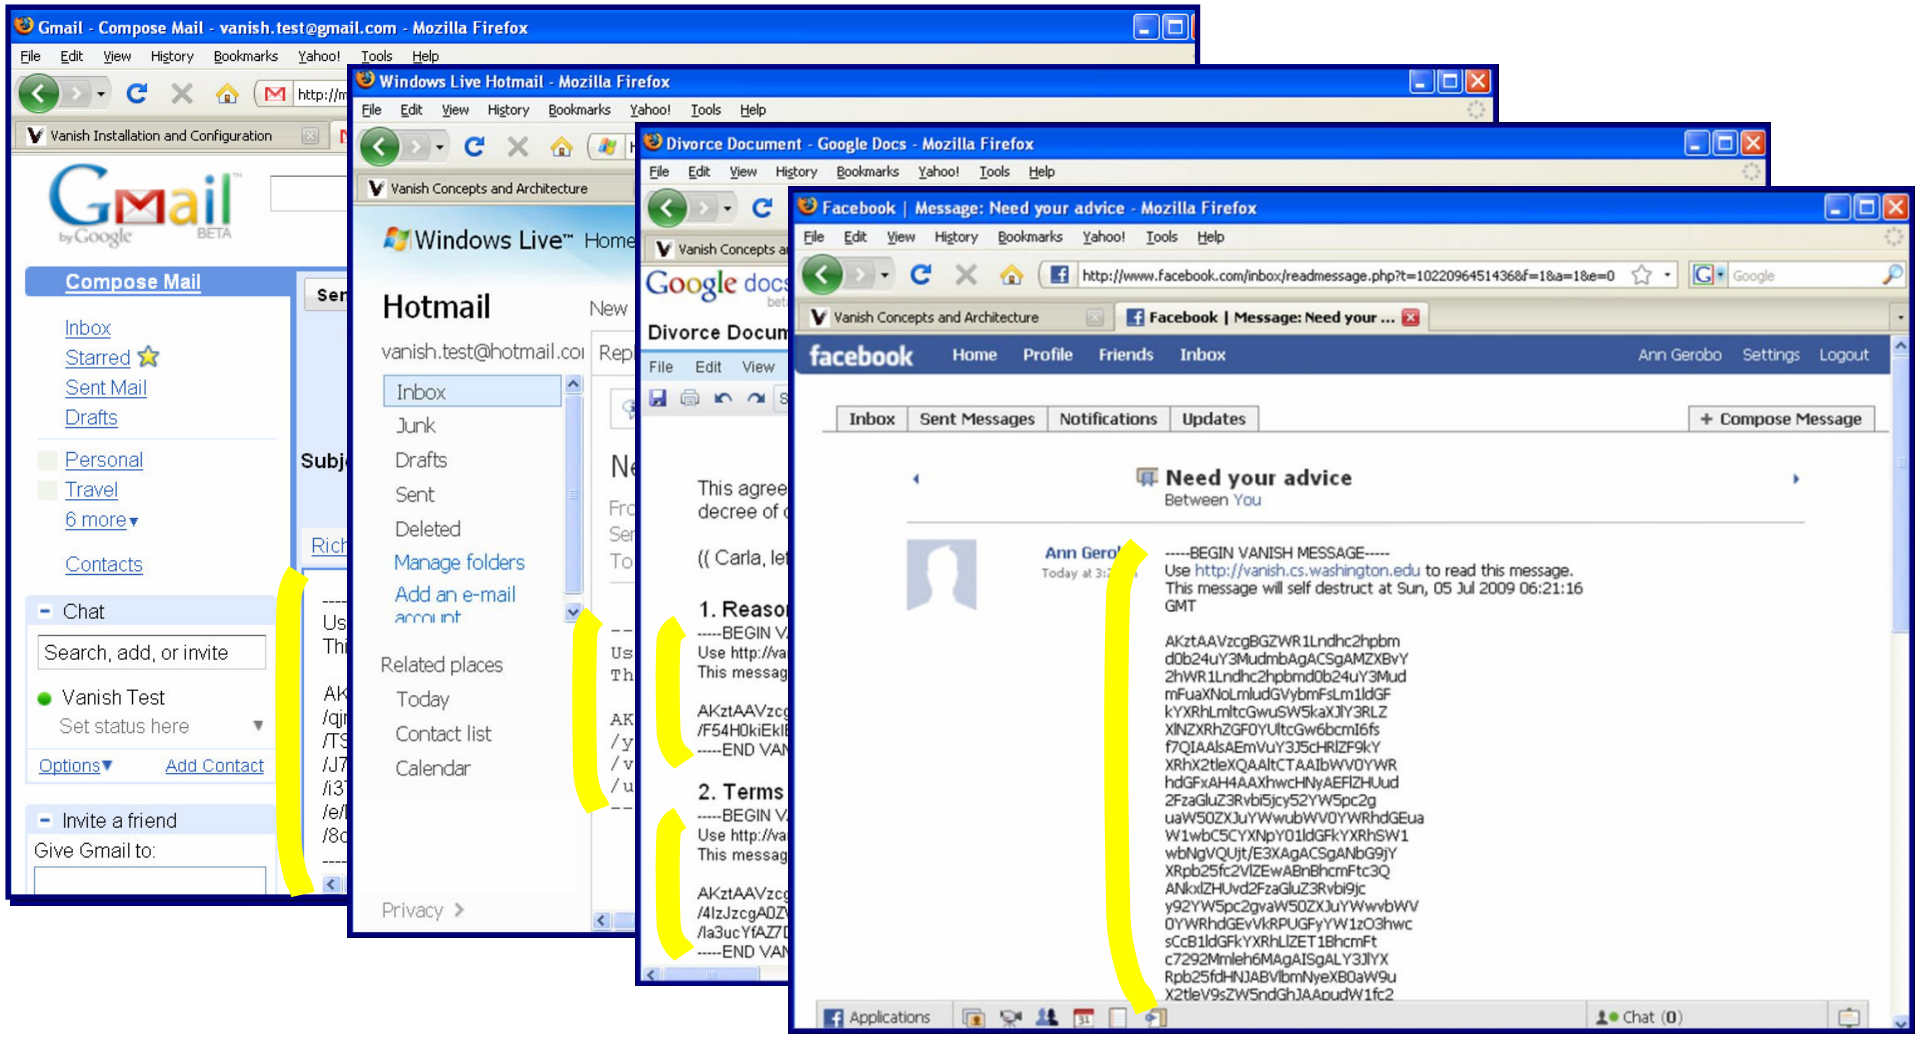

Firefox Plugin For Vanishing Web Data

- Encapsulate text in any text area in self-destructing VDOs

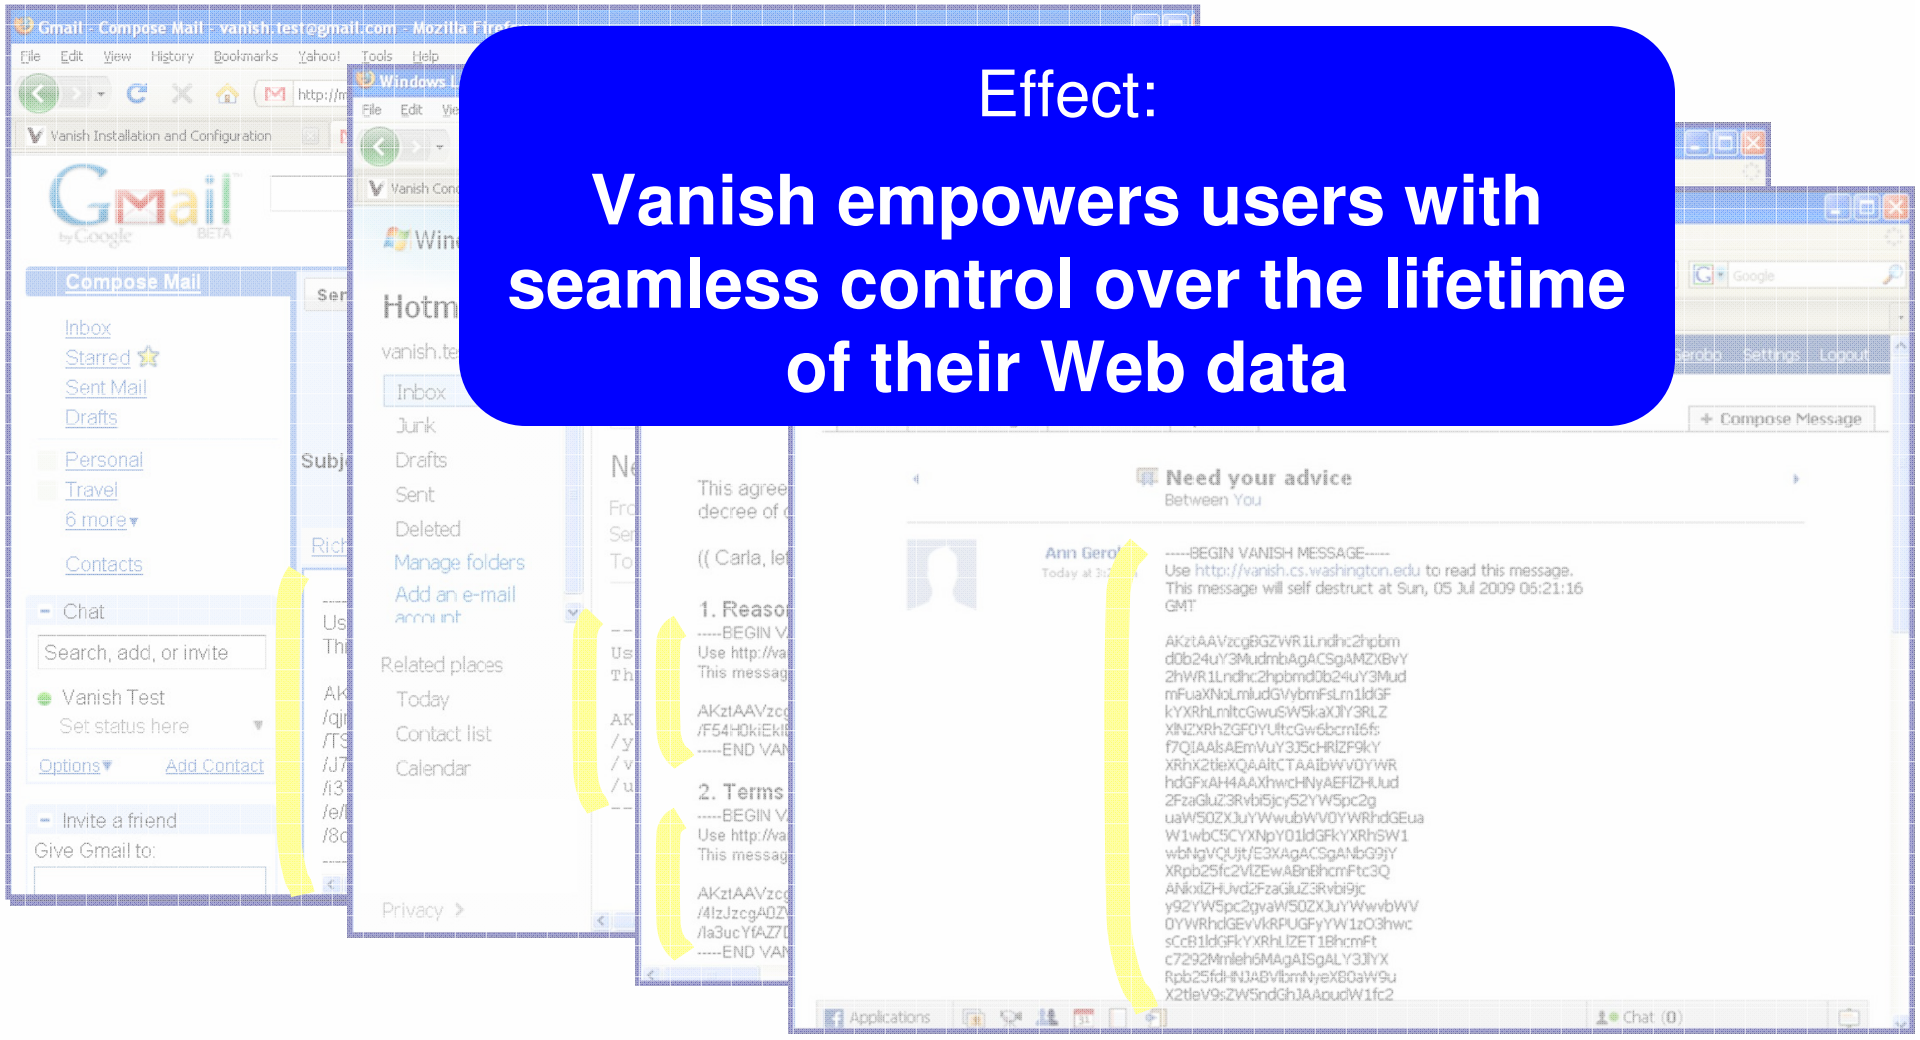

http://vanish.cs.washington.edu/

Conclusions

- Two formidable challenges to privacy:
  - Data lives forever
  - Disclosures of data and keys have become commonplace
- Self-destructing data empowers users with lifetime control
- Vanish:
  - Combines global-scale DHTs with secret sharing to provide self-destructing data
  - Firefox plugin allows users to set timeouts on text data anywhere on the web
- Vanish ≠ Vuze-based Vanish
  - Customized DHTs, hybrid approach, other P2P systems
  - Further extensions for security in the paper